

Chapter 04

Threads & Concurrency

A fundamental unit of CPU utilization

Outline

- ◆ Overview
- ◆ Multicore Programming
- ◆ Multithreading Models
- ◆ Thread Libraries
- ◆ Implicit Threading
- ◆ Threading Issues

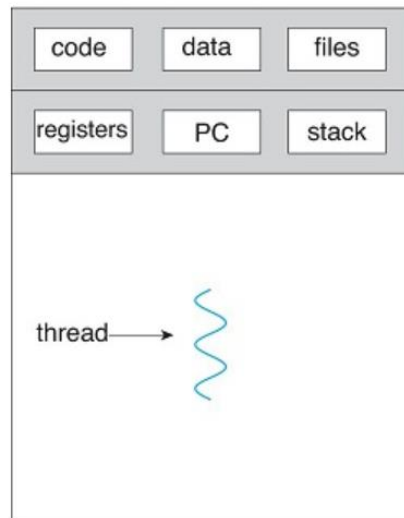
Objectives

- ◆ Identify the basic components of a thread, and contrast threads and processes.
- ◆ Describe the major benefits and significant challenges of designing multithreaded processes.
- ◆ Illustrate different approaches to implicit threading
- ◆ Describe threading issues


Overview

◆ Thread

- ❑ Basic unit of CPU utilization
- ❑ Include a **thread ID**, a **program counter**, a **register set**, and a **stack**
- ❑ Share code section, data section, and other resources such as open files and signals
- ❑ Program with multiple threads can perform multiple tasks



single-threaded process



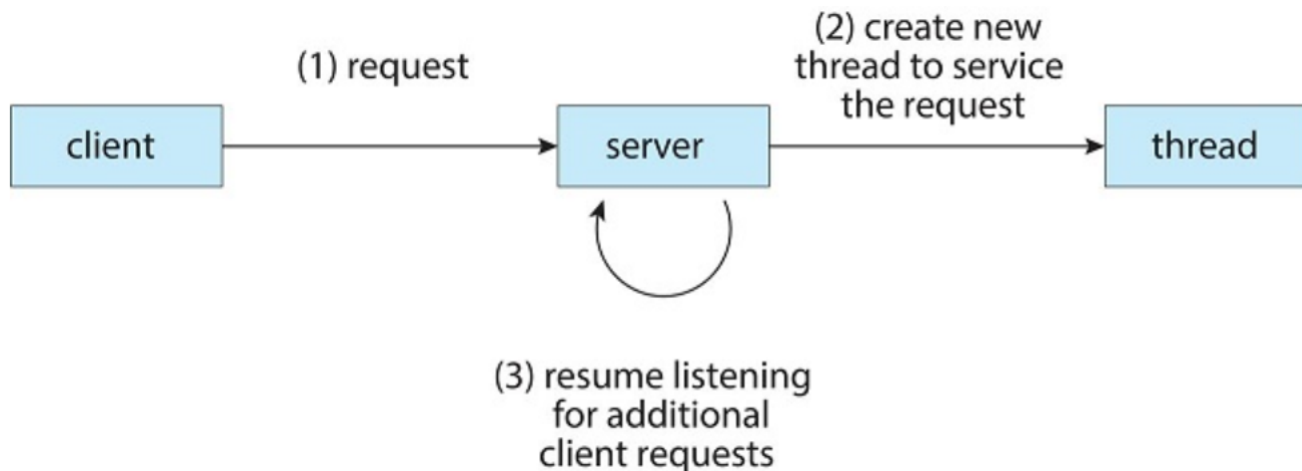
multithreaded process

Motivation

- ◆ Most modern applications are multithreaded
- ◆ Threads run within application
 - Thumbnails
 - Web browser
 - Word processor
- ◆ Leverage **processing capabilities** on multicore systems
- ◆ Consider a web server
 - Provide web pages, images, ...
 - Busy for multiple clients
 - Handle by traditional single-threaded process
 - One client at a time

Motivation

- ❑ Handle by multiple processes
 - Each process serves a request
 - Process is **heavy-weight**: time consuming and resource intensive
- ❑ Handle by on process with multiple threads
 - Each thread serves a request
 - Thread is **light-weight**



Motivation

- ◆ **Kernels** are generally multithreaded
 - ❑ Linux's kernel threads for varying tasks (shown by `ps -ef`)
 - ❑ `kthreadd` (with pid = 2) as parent of all kernel threads
- ◆ High performance computing applications via threads for running in **parallel**
 - ❑ Data mining
 - ❑ Graphics
 - ❑ Artificial intelligence

Benefits


- ◆ **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- ◆ **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- ◆ **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- ◆ **Scalability** – process can take advantage of multiprocessor architectures

Exercises


- ◆ Describe the actions taken by a kernel to context-switch between kernel-level threads.
- ◆ What resources are used when a thread is created? How do they differ from those used when a process is created?

Multicore Programming

- ◆ **Multicore** or **multiprocessor** systems bring parallelism and improve concurrency.
 - **Parallelism** implies a system can perform more than one task simultaneously
 - **Concurrency** supports more than one task making progress
 - Single processor / core, scheduler providing concurrency



Concurrent execution on a single-core system.



Parallel execution on a multicore system.

Exercise


- ◆ Is it possible to have concurrency but not parallelism?
Explain.

Programming Challenges

- ◆ The programming challenges in multicore programming include:
 - **Identifying tasks**
 - **Balance**
 - **Data splitting**
 - **Data dependency**
 - **Testing and debugging**

Types of parallelism

- ◆ **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
- ◆ **Task parallelism** – distributing threads across cores, each thread performing unique operation




- ◆ As # of threads grows, so does architectural support for threading
 - ❑ CPUs have cores as well as **hardware threads**
 - ❑ Consider Oracle SPARC T4 with 8 cores, and 8 hardware threads per core

Amdahl's Law

- ◆ Identifies performance gains from adding additional cores to an application that has both serial and parallel components
- ◆ Given S as serial portion and N processing cores, we have:

$$speedup \leq \frac{1}{S + \frac{1-S}{N}}$$

- ◆ As N approaches infinity, speedup approaches $1 / S$
- ◆ For instance, if application is 75% parallel / 25% serial
 - Moving from 1 to 2 cores results in speedup of 1.6 times
 - Moving to infinity cores results in speedup of 4 times




Serial portion of an application has disproportionate effect on performance gained by adding additional cores

Multithreading Models

- ◆ **User threads** - management done by user-level threads library. Three primary thread libraries:
 - ❑ POSIX **Pthreads**
 - ❑ Windows threads
 - ❑ Java threads
- ◆ **Kernel threads** - Supported by the Kernel. Examples include virtually all general purpose operating systems:
 - ❑ Windows
 - ❑ Solaris
 - ❑ Linux
 - ❑ Tru64 UNIX
 - ❑ Mac OS X

Multithreading Models


- ◆ Many-to-One
- ◆ One-to-One
- ◆ Many-to-Many



Many-to-One


- ◆ Many user-level threads mapped to single kernel thread
- ◆ One thread blocking causes all to block
- ◆ Multiple threads may not run in parallel on muticore system because only one may be in kernel at a time
- ◆ Few systems currently use this model
- ◆ Examples:

▣ Solaris Green Threads



One-to-One

- ◆ Each user-level thread maps to kernel thread
- ◆ Creating a user-level thread creates a kernel thread
- ◆ More concurrency than many-to-one
- ◆ Number of threads per process sometimes restricted due to overhead
- ◆ Examples
 - ❑ Windows
 - ❑ Linux




Exercise

- ◆ A system with two dual-core processors has four processors available for scheduling. A CPU-intensive application is running on this system. All input is performed at program start-up, when a single file must be opened. Similarly, all output is performed just before the program terminates, when the program results must be written to a single file. Between start-up and termination, the program is entirely CPU-bound. Your task is to improve the performance of this application by multithreading it. The application runs on a system that uses the one-to-one threading model (each user thread maps to a kernel thread).
 - ❑ How many threads will you create to perform the input and output? Explain.
 - ❑ How many threads will you create for the CPU-intensive portion of the application? Explain.

Many-to-Many Model

- ◆ Allows many user level threads to be mapped to many kernel threads
- ◆ Allows the operating system to create a sufficient number of kernel threads




Exercise

- ◆ Consider a multicore system and a multithreaded program written using the many-to-many threading model. Let the number of user-level threads in the program be greater than the number of processing cores in the system. Discuss the performance implications of the following scenarios.
 - ❑ The number of kernel threads allocated to the program is less than the number of processing cores.
 - ❑ The number of kernel threads allocated to the program is equal to the number of processing cores.
 - ❑ The number of kernel threads allocated to the program is greater than the number of processing cores but less than the number of user-level threads.

Two-level Model

- ◆ Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- ◆ Examples
 - IRIX
 - HP-UX
 - Tru64 UNIX
 - Solaris 8 and earlier



Comparisons among Models

- ◆ Many-to-one
 - The number of user threads does not imply parallelism
- ◆ One-to-one
 - Limited number of threads can be used.
- ◆ Many-to-many
 - Flexible
 - Model suffers from neither of these shortcomings.
- ◆ Most operating systems use one-to-one model
 - Many-to-many is difficult to implement.
 - The increasing number of cores lowers down the importance of the number of kernel threads.

Exercises

- ◆ Provide a programming example in which multithreading does not provide better performance than a single-threaded solution.
- ◆ Can a multithreaded solution using multiple user-level threads achieve better performance on a multiprocessor system than on a single-processor system? Explain.

Thread Libraries

- ◆ Thread library provides programmer with API for creating and managing threads
- ◆ Two primary ways of implementing
 - ❑ Library entirely in user space
 - ❑ Kernel-level library supported by the OS
- ◆ Three main thread libraries
 - ❑ POSIX Pthreads
 - Either user level or kernel level
 - ❑ Windows
 - Kernel level
 - ❑ Java
 - Managed in Java program.
 - Implemented using a thread library available on the host system since JVM is running on top of a host operating system.

Thread Libraries

◆ General strategies for creating multiple threads

□ Asynchronous threading

- Parent resumes its execution after generating child thread
- Run parent and child concurrently and independently
- Little data sharing

□ Synchronous threading

- Parent waits for all of its children threads to terminate
- Children threads run concurrently
- Significant data sharing

Exercise

- ◆ Consider the following code segment:

```
pid_t pid;  
pid = fork();  
if (pid == 0) { /* child process */  
    fork();  
    thread_create( . . . );  
}  
fork();
```

- How many unique processes are created?
- How many unique threads are created?

Implicit Threading

- ◆ Growing in popularity as numbers of threads increase, program correctness more difficult with explicit threads
- ◆ Creation and management of threads done by compilers and run-time libraries rather than programmers. This is known as implicit threading.
 - Identify tasks (function) rather than threads which can run in parallel.
 - Map task to a separate thread via M-M model.
 - Developers only need to identify parallel tasks.
 - Libraries determine thread creation and management.
- ◆ Two methods explored
 - Thread Pools
 - OpenMP

Thread Pools

- ◆ Recall the multithreaded web server with two issues:
 - ❑ The amount of time required to create the thread
 - ❑ No bound on the number of active threads
- ◆ A solution is to use a thread pool.
 - ❑ Create a number of threads in a pool where they await work
 - ❑ A request is served by awaking a thread in pool.
 - ❑ Requests are queued if no thread is available.
 - ❑ A thread returns to the pool when task has completed.
 - ❑ Work well for asynchronous execution

Thread Pools

◆ Advantages of using thread pools

- ❑ Usually slightly faster to service a request with an existing thread than create a new thread
- ❑ Allows the number of threads in the application(s) to be bound to the size of the pool
- ❑ Separating task to be performed from mechanics of creating task allows different strategies for running task
 - i.e. tasks could be scheduled to run periodically

◆ The number of threads in the pool

- ❑ Can be set heuristically such as according to #CPUs, size of main memory, #concurrent requests
- ❑ Can be set dynamically according to usage pattern
 - E.g. Apple's Grand Central Dispatch

OpenMP

- ◆ Set of **compiler directives** and an API for C, C++, FORTRAN
- ◆ Provides support for parallel programming in shared-memory environments
- ◆ Identifies parallel regions – blocks of code that can run in parallel
- ◆ **#pragma omp parallel**
 - ❑ Create as many threads as there are cores

```
1  #include <omp.h>
2  #include <stdio.h>
3
4  int main(int argc, char *argv[])
5  {
6      /* sequential code */
7
8      #pragma omp parallel
9      {
10         printf("I am a parallel region\n");
11     }
12
13     /* sequential code */
14
15     return 0;
16 }
```

OpenMP

◆ Parallelizing loops

```
□ #pragma omp parallel for
    for(int i = 0; i < N; i++) {
        c[i] = a[i] + b[i];
    }
```

- Create as many threads as there are cores
- Run for loop in parallel by using
#pragma omp parallel for

Threading Issues

- ◆ Semantics of **fork()** and **exec()** system calls
- ◆ Signal handling
 - Synchronous and asynchronous
- ◆ Thread cancellation of target thread
 - Asynchronous or deferred
- ◆ Thread-local storage
- ◆ Scheduler Activations

Semantics of `fork()` and `exec()`

- ◆ Does `fork()` duplicate **only the calling thread** or **all threads**?
 - Some UNIXes have two versions of `fork()`
- ◆ `exec()` usually works as normal – replace the running process including all threads
- ◆ Depends on application
 - `fork()` then `exec()`: Duplicate calling thread is okay.
 - `fork()` without `exec()`: Duplicate all thread is needed.

Signal Handling

- ◆ **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- ◆ A **signal handler** is used to process signals
 - ❑ Signal is generated by particular event
 - ❑ Signal is delivered to a process
 - ❑ Signal is handled after delivered
- ◆ **Synchronous signals**
 - ❑ Delivered to the same process that performed the operation that caused the signal
 - ❑ E.g. illegal memory access and division by 0
- ◆ **Asynchronous signals**
 - ❑ Generated by an external event
 - ❑ Sent to another process
 - ❑ E.g. terminating a process with specific keystrokes, timer expire

Signal Handling

- ◆ A signal may be ***handled*** by one of two possible handlers:
 - default
 - user-defined
- ◆ Every signal has **default handler** that kernel runs when handling signal
 - **User-defined signal handler** can override default

Signal Handling

- ◆ Where should a signal be delivered for multi-threaded?
 - ❑ Deliver the signal to the thread to which the signal applies
 - ❑ Deliver the signal to every thread in the process
 - ❑ Deliver the signal to certain threads in the process
 - ❑ Assign a specific thread to receive all signals for the process

- ◆ The method for delivering a signal depends on the type of signal generated.
 - ❑ Synchronous signal to the thread causing the signal.
 - ❑ Asynchronous signal depends.

Thread Cancellation

- ◆ Terminating a thread before it has finished
- ◆ Thread to be canceled is **target thread**
- ◆ Two general approaches:
 - ❑ **Asynchronous cancellation** terminates the target thread immediately
 - ❑ **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
- ◆ Difficulty with cancellation
 - ❑ Allocated resources
 - ❑ In the midst of updating shared data
 - ❑ Troublesome with asynchronous cancellation
 - May not free all resources
 - ❑ Safely canceled by deferred cancellation

Thread Cancellation

- ◆ Invoking thread cancellation requests cancellation, but actual cancellation depends on thread state


Mode	State	Type
Off	Disabled	–
Deferred	Enabled	Deferred
Asynchronous	Enabled	Asynchronous

- ◆ If thread has cancellation disabled, cancellation remains pending until thread enables it
- ◆ Default type is deferred
 - ❑ Cancellation only occurs when thread reaches **cancellation point**
 - ❑ Then **cleanup handler** is invoked
- ◆ On Linux systems, thread cancellation is handled through signals

Thread-Local Storage

- ◆ Thread-local storage (TLS) allows each thread to have its own copy of data
- ◆ Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- ◆ Different from local variables
 - ❑ Local variables visible only during single function invocation
 - ❑ TLS visible across function invocations
- ◆ Similar to **static** data
 - ❑ TLS is unique to each thread.
 - ❑ Most thread libraries and compilers provide support for TLS.

Scheduler Activations



- ◆ Both M:M and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- ◆ Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)**
 - Appears to be a virtual processor on which process can schedule user thread to run
 - Each LWP attached to kernel thread
- ◆ Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the **upcall handler** in the thread library
- ◆ This communication allows an application to maintain the correct number kernel threads

Exercise

- ◆ Assume that an operating system maps user-level threads to the kernel using the many-to-many model and that the mapping is done through LWPs. Furthermore, the system allows developers to create real-time threads for use in real-time systems. Is it necessary to bind a real-time thread to an LWP? Explain.