


# Chapter 6: Process Scheduling

---





# Chapter 6: Process Scheduling

---

- Background
- The Critical-Section Problem
- Peterson's Solution
- Synchronization Hardware
- Mutex Locks
- Semaphores
- Classic Problems of Synchronization
- Monitors
- Synchronization Examples
- Alternative Approaches





# Objectives

---

To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data

To present both software and hardware solutions of the critical-section problem

To examine several classical process-synchronization problems

To explore several tools that are used to solve process synchronization problems





# Background

---

Processes can execute concurrently

May be interrupted at any time, partially completing execution

Concurrent access to shared data may result in data inconsistency

Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes

Illustration of the problem:

Suppose that we wanted to provide a solution to the consumer-producer problem that fills **all** the buffers. We can do so by having an integer **counter** that keeps track of the number of full buffers. Initially, **counter** is set to 0. It is incremented by the producer after it produces a new buffer and is decremented by the consumer after it consumes a buffer.





# Producer

---

```
while (true) {  
    /* produce an item in next produced */  
  
    while (counter == BUFFER SIZE) ;  
        /* do nothing */  
    buffer[in] = next produced;  
    in = (in + 1) % BUFFER SIZE;  
    counter++;  
}
```





# Consumer

```
while (true) {  
    while (counter == 0)  
        ; /* do nothing */  
    next consumed = buffer[out];  
    out = (out + 1) % BUFFER SIZE;      counter-  
-;  
    /* consume the item in next consumed */  
}
```





# Race Condition

**counter++** could be implemented as

```
register1 = counter
register1 = register1 + 1
counter = register1
```

**counter--** could be implemented as

```
register2 = counter
register2 = register2 - 1
counter = register2
```

Consider this execution interleaving with “count = 5” initially:

S0: producer execute	<b>register1 = counter</b>	{register1 = 5}
S1: producer execute	<b>register1 = register1 + 1</b>	{register1 = 6}
S2: consumer execute	<b>register2 = counter</b>	{register2 = 5}
S3: consumer execute	<b>register2 = register2 - 1</b>	{register2 = 4}
S4: producer execute	<b>counter = register1</b>	{counter = 6}
S5: consumer execute	<b>counter = register2</b>	{counter = 4}





# Critical Section Problem

---

Consider system of  $n$  processes  $\{p_0, p_1, \dots, p_{n-1}\}$

Each process has **critical section** segment of code

Process may be changing common variables, updating table, writing file, etc


When one process in critical section, no other may be in its critical section

**Critical section problem** is to design protocol to solve this

Each process must ask permission to enter critical section in **entry section**, may follow critical section with **exit section**, then **remainder section**







# Critical Section

General structure of process  $p_i$  is

```
do {  
    entry section  
    critical section  
    exit section  
    remainder section  
} while (true);
```





# Solution to Critical-Section Problem

1. **Mutual Exclusion** - If process  $P_i$  is executing in its critical section, then no other processes can be executing in their critical sections
2. **Progress** - If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely
3. **Bounded Waiting** - A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted
  - Assume that each process executes at a nonzero speed
  - No assumption concerning **relative speed** of the  $n$  processes

Two approaches depending on if kernel is preemptive or non-preemptive

**Preemptive** – allows preemption of process when running in kernel mode

**Non-preemptive** – runs until exits kernel mode, blocks, or voluntarily yields CPU

▶ Essentially free of race conditions in kernel mode





# Peterson's Solution

---

Good algorithmic description of solving the problem

Two process solution

Assume that the **load** and **store** instructions are atomic; that is, cannot be interrupted

The two processes share two variables:

```
int turn;
```

```
Boolean flag[2]
```

The variable **turn** indicates whose turn it is to enter the critical section

The **flag** array is used to indicate if a process is ready to enter the critical section. **flag[i] = true** implies that process  $P_i$  is ready!





# Algorithm for Process $P_i$

```
do {  
    flag[i] = true;  
    turn = j;  
    while (flag[j] && turn == j);  
        critical section  
    flag[i] = false;  
        remainder section  
} while (true);
```

Provable that

1. Mutual exclusion is preserved
2. Progress requirement is satisfied
3. Bounded-waiting requirement is met





# Synchronization Hardware

---

Many systems provide hardware support for critical section code

All solutions below based on idea of **locking**

Protecting critical regions via locks

Uniprocessors – could disable interrupts

Currently running code would execute without preemption

Generally too inefficient on multiprocessor systems

- ▶ Operating systems using this not broadly scalable

Modern machines provide special atomic hardware instructions

- ▶ **Atomic** = non-interruptible

Either test memory word and set value

Or swap contents of two memory words





# Solution to Critical-section Problem Using Locks

```
do {  
    acquire lock  
        critical section  
    release lock  
        remainder section  
} while (TRUE);
```





# test\_and\_set Instruction

---

Definition:

```
boolean test_and_set (boolean *target)
{
    boolean rv = *target;
    *target = TRUE;
    return rv;
}
```






# Solution using test\_and\_set()

---

Shared boolean variable lock, initialized to FALSE

Solution:

```
do {  
    while (test_and_set(&lock))  
        ; /* do nothing */  
    /* critical section */  
    lock = false;  
    /* remainder section */  
} while (true);
```





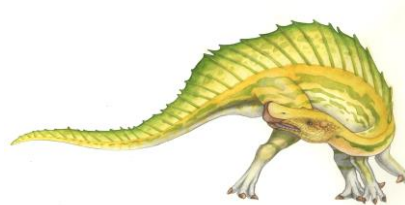


# compare\_and\_swap Instruction

---

Definition:

```
int compare_and_swap(int *value, int expected, int new value) {  
    int temp = *value;  
    if (*value == expected)  
        *value = new value;  
    return temp;  
}
```






# Solution using compare\_and\_swap

Shared Boolean variable lock initialized to FALSE; Each process has a local Boolean variable key  
Solution:

```
do {  
    while (compare and swap(&lock, 0, 1) != 0)  
        ; /* do nothing */  
    /* critical section */  
    lock = 0;  
    /* remainder section */  
} while (true);
```





# Bounded-waiting Mutual Exclusion with test\_and\_set

```
do {
    waiting[i] = true;
    key = true;
    while (waiting[i] && key)
        key = test_and_set(&lock);
    waiting[i] = false;
    /* critical section */
    j = (i + 1) % n;
    while ((j != i) && !waiting[j])
        j = (j + 1) % n;
    if (j == i)
        lock = false;
    else
        waiting[j] = false;
    /* remainder section */
} while (true);
```





# Mutex Locks

Previous solutions are complicated and generally inaccessible to application programmers

OS designers build software tools to solve critical section problem

Simplest is mutex lock

Protect critical regions with it by first **acquire()** a lock then **release()** it

Boolean variable indicating if lock is available or not

Calls to **acquire()** and **release()** must be atomic

Usually implemented via hardware atomic instructions

But this solution requires **busy waiting**

This lock therefore called a **spinlock**





# acquire() and release()

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;;  
}  
release() {  
    available = true;  
}  
  
do {  
    acquire lock  
    critical section  
    release lock  
    remainder section  
} while (true);
```





# Semaphore

Synchronization tool that does not require busy waiting

Semaphore **S** – integer variable

Two standard operations modify **S**: **wait()** and **signal()**

Originally called **P()** and **V()**

Less complicated

Can only be accessed via two indivisible (atomic) operations

```
wait (S) {  
    while (S <= 0)  
        ; // busy wait  
    S--;  
}  
signal (S) {  
    S++;  
}
```





# Semaphore Usage

**Counting semaphore** – integer value can range over an unrestricted domain

**Binary semaphore** – integer value can range only between 0 and 1

Then a **mutex lock**

Can implement a counting semaphore **S** as a binary semaphore

Can solve various synchronization problems

Consider **P<sub>1</sub>** and **P<sub>2</sub>** that require **S<sub>1</sub>** to happen before **S<sub>2</sub>**

P1 :

**S<sub>1</sub>** ;

**signal (synch)** ;

P2 :

**wait (synch)** ;

**S<sub>2</sub>** ;





# Semaphore Implementation

---

Must guarantee that no two processes can execute `wait ()` and `signal ()` on the same semaphore at the same time

Thus, implementation becomes the critical section problem where the wait and signal code are placed in the critical section

Could now have **busy waiting** in critical section implementation

- ▶ But implementation code is short
- ▶ Little busy waiting if critical section rarely occupied

Note that applications may spend lots of time in critical sections and therefore this is not a good solution







# Semaphore Implementation with no Busy waiting

---

With each semaphore there is an associated waiting queue

Each entry in a waiting queue has two data items:

- value (of type integer)

- pointer to next record in the list

Two operations:

- block** – place the process invoking the operation on the appropriate waiting queue

- wakeup** – remove one of processes in the waiting queue and place it in the ready queue





# Semaphore Implementation with no Busy waiting (Cont.)

```
typedef struct{
    int value;
    struct process *list;
} semaphore;

wait(semaphore *S) {
    S->value--;
    if (S->value < 0) {
        add this process to S->list;
        block();
    }
}

signal(semaphore *S) {
    S->value++;
    if (S->value <= 0) {
        remove a process P from S->list;
        wakeup(P);
    }
}
```





# Deadlock and Starvation

**Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes

Let  $S$  and  $Q$  be two semaphores initialized to 1

$P_0$   
`wait(S) ;`  
`wait(Q) ;`  
  
.  
`signal(S) ;`  
`signal(Q) ;`

$P_1$   
`wait(Q) ;`  
`wait(S) ;`  
  
.  
`signal(Q) ;`  
`signal(S) ;`

**Starvation** – indefinite blocking

A process may never be removed from the semaphore queue in which it is suspended

**Priority Inversion** – Scheduling problem when lower-priority process holds a lock needed by higher-priority process

Solved via **priority-inheritance protocol**





# Classical Problems of Synchronization

---

Classical problems used to test newly-proposed synchronization schemes

Bounded-Buffer Problem

Readers and Writers Problem

Dining-Philosophers Problem





# Bounded-Buffer Problem

---


$n$  buffers, each can hold one item

Semaphore `mutex` initialized to the value 1

Semaphore `full` initialized to the value 0

Semaphore `empty` initialized to the value  $n$





# Bounded Buffer Problem (Cont.)

The structure of the producer process

```
do {  
    ...  
    /* produce an item in next_produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    ...  
    /* add next produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while (true);
```





# Bounded Buffer Problem (Cont.)

The structure of the consumer process

```
do {  
    wait(full);  
    wait(mutex);  
    ...  
    /* remove an item from buffer to next_consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
    ...  
    /* consume the item in next consumed */  
    ...  
} while (true);
```





# Readers-Writers Problem

---

A data set is shared among a number of concurrent processes

Readers – only read the data set; they do **not** perform any updates

Writers – can both read and write

Problem – allow multiple readers to read at the same time

Only one single writer can access the shared data at the same time

Several variations of how readers and writers are treated – all involve priorities

Shared Data

Data set

Semaphore **rw\_mutex** initialized to 1

Semaphore **mutex** initialized to 1

Integer **read\_count** initialized to 0





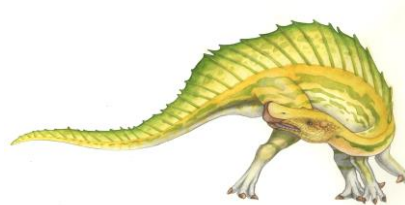


# Readers-Writers Problem (Cont.)

---

The structure of a writer process

```
do {  
    wait(rw mutex);  
    ...  
    /* writing is performed */  
    ...  
    signal(rw mutex);  
} while (true);
```






# Readers-Writers Problem (Cont.)

The structure of a reader process

```
do {  
    wait(mutex);  
    read count++;  
    if (read count == 1)  
        wait(rw mutex);  
    signal(mutex);  
  
    ...  
    /* reading is performed */  
    ... wait(mutex);  
    read count--;  
    if (read count == 0)  
        signal(rw mutex);  
    signal(mutex);  
} while (true);
```





# Readers-Writers Problem Variations

---

*First* variation – no reader kept waiting unless writer has permission to use shared object

*Second* variation – once writer is ready, it performs write asap


Both may have starvation leading to even more variations

Problem is solved on some systems by kernel providing reader-writer locks





# Dining-Philosophers Problem



Philosophers spend their lives thinking and eating

Don't interact with their neighbors, occasionally try to pick up 2 chopsticks (one at a time) to eat from bowl

Need both to eat, then release both when done

In the case of 5 philosophers

Shared data

- ▶ Bowl of rice (data set)
- ▶ Semaphore **chopstick** [5] initialized to 1





# Dining-Philosophers Problem Algorithm

The structure of Philosopher  $i$ :

```
do {  
    wait ( chopstick[i] );  
    wait ( chopStick[ (i + 1) % 5] );  
  
    // eat  
  
    signal ( chopstick[i] );  
    signal ( chopstick[ (i + 1) % 5] );  
  
    // think  
  
} while (TRUE);
```

What is the problem with this algorithm?





# Problems with Semaphores

---

Incorrect use of semaphore operations:

signal (mutex) .... wait (mutex)

wait (mutex) ... wait (mutex)

Omitting of wait (mutex) or signal (mutex) (or both)

Deadlock and starvation





# Monitors

A high-level abstraction that provides a convenient and effective mechanism for process synchronization

*Abstract data type*, internal variables only accessible by code within the procedure


Only one process may be active within the monitor at a time

But not powerful enough to model some synchronization schemes

```
monitor monitor-name
{
    // shared variable declarations
    procedure P1 (...) { .... }


    procedure Pn (...) {.....}

    Initialization code (...) { ... }
}
}
```





# Schematic view of a Monitor







# Condition Variables

---

condition x, y;

Two operations on a condition variable:

`x.wait ()` – a process that invokes the operation is suspended until `x.signal ()`


`x.signal ()` – resumes one of processes (if any) that invoked `x.wait ()`

- ▶ If no `x.wait ()` on the variable, then it has no effect on the variable





# Monitor with Condition Variables





# Condition Variables Choices

If process P invokes `x.signal ()`, with Q in `x.wait ()` state, what should happen next?

If Q is resumed, then P must wait

Options include

**Signal and wait** – P waits until Q leaves monitor or waits for another condition

**Signal and continue** – Q waits until P leaves the monitor or waits for another condition

Both have pros and cons – language implementer can decide

Monitors implemented in Concurrent Pascal compromise

- ▶ P executing signal immediately leaves the monitor, Q is resumed

Implemented in other languages including Mesa, C#, Java





# Solution to Dining Philosophers

---

monitor DiningPhilosophers

{

enum { THINKING; HUNGRY, EATING) state [5] ;  
condition self [5];

void pickup (int i) {  
    state[i] = HUNGRY;  
    test(i);  
    if (state[i] != EATING) self [i].wait;  
}

void putdown (int i) {  
    state[i] = THINKING;  
    // test left and right neighbors  
    test((i + 4) % 5);  
    test((i + 1) % 5);  
}





# Solution to Dining Philosophers (Cont.)

```
void test (int i) {  
    if ( (state[(i + 4) % 5] != EATING) &&  
        (state[i] == HUNGRY) &&  
        (state[(i + 1) % 5] != EATING) ) {  
        state[i] = EATING ;  
        self[i].signal () ;  
    }  
}
```

```
initialization_code() {  
    for (int i = 0; i < 5; i++)  
        state[i] = THINKING;  
}
```





# Solution to Dining Philosophers (Cont.)

---

Each philosopher  $i$  invokes the operations `pickup()` and `putdown()` in the following sequence:

`DiningPhilosophers.pickup (i);`

EAT

`DiningPhilosophers.putdown (i);`

No deadlock, but starvation is possible





# Monitor Implementation Using Semaphores


Variables

```
semaphore mutex; // (initially = 1)
semaphore next;  // (initially = 0)
int next_count = 0;
```

Each procedure  $F$  will be replaced by

```
wait(mutex);
...
    body of  $F$ ;
...
if (next_count > 0)
    signal(next)
else
    signal(mutex);
```

Mutual exclusion within a monitor is ensured





# Monitor Implementation – Condition Variables

For each condition variable  $x$ , we have:

```
semaphore x_sem; // (initially = 0)
int x_count = 0;
```

The operation  $x.wait$  can be implemented as:

```
x-count++;
if (next_count > 0)
    signal(next);
else
    signal(mutex);
wait(x_sem);
x-count--;
```








# Monitor Implementation (Cont.)

---

The operation `x.signal` can be implemented as:

```
if (x-count > 0) {  
    next_count++;  
    signal(x_sem);  
    wait(next);  
    next_count--;  
}
```





# Resuming Processes within a Monitor

If several processes queued on condition x, and x.signal() executed, which should be resumed?

FCFS frequently not adequate

**conditional-wait** construct of the form x.wait(c)

Where c is **priority number**

Process with lowest number (highest priority) is scheduled next






# A Monitor to Allocate Single Resource

```
monitor ResourceAllocator
{
    boolean busy;
    condition x;
    void acquire(int time) {
        if (busy)
            x.wait(time);
        busy = TRUE;
    }
    void release() {
        busy = FALSE;
        x.signal();
    }
    initialization code() {
        busy = FALSE;
    }
}
```





# Synchronization Examples

---

Solaris

Windows XP

Linux

Pthreads





# Solaris Synchronization

---

Implements a variety of locks to support multitasking, multithreading (including real-time threads), and multiprocessing

Uses **adaptive mutexes** for efficiency when protecting data from short code segments

- Starts as a standard semaphore spin-lock

- If lock held, and by a thread running on another CPU, spins

- If lock held by non-run-state thread, block and sleep waiting for signal of lock being released

Uses **condition variables**

Uses **readers-writers** locks when longer sections of code need access to data

Uses **turnstiles** to order the list of threads waiting to acquire either an adaptive mutex or reader-writer lock

- Turnstiles are per-lock-holding-thread, not per-object

Priority-inheritance per-turnstile gives the running thread the highest of the priorities of the threads in its turnstile





# Windows XP Synchronization

---

Uses interrupt masks to protect access to global resources on uniprocessor systems

Uses **spinlocks** on multiprocessor systems

Spinlocking-thread will never be preempted

Also provides **dispatcher objects** user-land which may act mutexes, semaphores, events, and timers

## Events

- ▶ An event acts much like a condition variable

Timers notify one or more thread when time expired

Dispatcher objects either **signaled-state** (object available) or **non-signaled state** (thread will block)





# Linux Synchronization

---

Linux:

Prior to kernel Version 2.6, disables interrupts to implement short critical sections

Version 2.6 and later, fully preemptive

Linux provides:

semaphores

spinlocks

reader-writer versions of both

On single-cpu system, spinlocks replaced by enabling and disabling kernel preemption





# Pthreads Synchronization

---

Pthreads API is OS-independent

It provides:


- mutex locks

- condition variables


Non-portable extensions include:

- read-write locks

- spinlocks







# Atomic Transactions

---

System Model

Log-based Recovery

Checkpoints

Concurrent Atomic Transactions





# System Model

---

Assures that operations happen as a single logical unit of work, in its entirety, or not at all

Related to field of database systems

Challenge is assuring atomicity despite computer system failures

**Transaction** - collection of instructions or operations that performs single logical function

Here we are concerned with changes to stable storage – disk

Transaction is series of **read** and **write** operations

Terminated by **commit** (transaction successful) or **abort** (transaction failed) operation

Aborted transaction must be **rolled back** to undo any changes it performed





# Types of Storage Media

---

Volatile storage – information stored here does not survive system crashes

Example: main memory, cache


Nonvolatile storage – Information usually survives crashes

Example: disk and tape

Stable storage – Information never lost

Not actually possible, so approximated via replication or RAID to devices with independent failure modes

Goal is to assure transaction atomicity where failures cause loss of information on volatile storage





# Log-Based Recovery

Record to stable storage information about all modifications by a transaction

Most common is [write-ahead logging](#)

Log on stable storage, each log record describes single transaction write operation, including

- ▶ Transaction name
- ▶ Data item name
- ▶ Old value
- ▶ New value

< $T_i$  starts> written to log when transaction  $T_i$  starts

< $T_i$  commits> written when  $T_i$  commits

Log entry must reach stable storage before operation on data occurs





# Log-Based Recovery Algorithm

---

Using the log, system can handle any volatile memory errors

**Undo( $T_i$ )** restores value of all data updated by  $T_i$

**Redo( $T_i$ )** sets values of all data in transaction  $T_i$  to new values

Undo( $T_i$ ) and redo( $T_i$ ) must be **idempotent**


Multiple executions must have the same result as one execution

If system fails, restore state of all updated data via log

If log contains  $\langle T_i \text{ starts} \rangle$  without  $\langle T_i \text{ commits} \rangle$ , **undo( $T_i$ )**

If log contains  $\langle T_i \text{ starts} \rangle$  and  $\langle T_i \text{ commits} \rangle$ , **redo( $T_i$ )**





# Checkpoints

---

Log could become long, and recovery could take long

Checkpoints shorten log and recovery time.

Checkpoint scheme:

1. Output all log records currently in volatile storage to stable storage
2. Output all modified data from volatile to stable storage
3. Output a log record <checkpoint> to the log on stable storage

Now recovery only includes Ti, such that Ti started executing before the most recent checkpoint, and all transactions after Ti. All other transactions already on stable storage.





# Concurrent Transactions

---


Must be equivalent to serial execution – [serializability](#)

Could perform all transactions in critical section

Inefficient, too restrictive

[Concurrency-control algorithms](#) provide serializability





# Serializability

---

Consider two data items A and B

Consider Transactions  $T_0$  and  $T_1$

Execute  $T_0$ ,  $T_1$  atomically

Execution sequence called **schedule**

Atomically executed transaction order called **serial schedule**

For N transactions, there are  $N!$  valid serial schedules







# Schedule 1: $T_0$ then $T_1$

$T_0$	$T_1$
read( $A$ )	
write( $A$ )	
read( $B$ )	
write( $B$ )	
	read( $A$ )
	write( $A$ )
	read( $B$ )
	write( $B$ )





# Nonserial Schedule

Nonserial schedule allows overlapped execute

Resulting execution not necessarily incorrect

Consider schedule  $S$ , operations  $O_i, O_j$

**Conflict** if access same data item, with at least one write

If  $O_i, O_j$  consecutive and operations of different transactions &  $O_i$  and  $O_j$  don't conflict

Then  $S'$  with swapped order  $O_j O_i$  equivalent to  $S$

If  $S$  can become  $S'$  via swapping nonconflicting operations

$S$  is **conflict serializable**





# Schedule 2: Concurrent Serializable Schedule

$T_0$	$T_1$
read( $A$ ) write( $A$ )	read( $A$ ) write( $A$ )
read( $B$ ) write( $B$ )	read( $B$ ) write( $B$ )





# Locking Protocol

---

Ensure serializability by associating lock with each data item

Follow locking protocol for access control

Locks

**Shared** –  $T_i$  has shared-mode lock (S) on item Q,  $T_i$  can read Q but not write Q

**Exclusive** –  $T_i$  has exclusive-mode lock (X) on Q,  $T_i$  can read and write Q

Require every transaction on item Q acquire appropriate lock

If lock already held, new request may have to wait

Similar to readers-writers algorithm





# Two-phase Locking Protocol

---

Generally ensures conflict serializability

Each transaction issues lock and unlock requests in two phases

- Growing – obtaining locks

- Shrinking – releasing locks

Does not prevent deadlock





# Timestamp-based Protocols

---

Select order among transactions in advance – [timestamp-ordering](#)

Transaction  $T_i$  associated with timestamp  $TS(T_i)$  before  $T_i$  starts

$TS(T_i) < TS(T_j)$  if  $T_i$  entered system before  $T_j$

TS can be generated from system clock or as logical counter incremented at each entry of transaction

Timestamps determine serializability order

If  $TS(T_i) < TS(T_j)$ , system must ensure produced schedule equivalent to serial schedule where  $T_i$  appears before  $T_j$





# Timestamp-based Protocol Implementation

Data item Q gets two timestamps

W-timestamp(Q) – largest timestamp of any transaction that executed write(Q) successfully

R-timestamp(Q) – largest timestamp of successful read(Q)

Updated whenever read(Q) or write(Q) executed

Timestamp-ordering protocol assures any conflicting read and write executed in timestamp order

Suppose  $T_i$  executes read(Q)

If  $TS(T_i) < W\text{-timestamp}(Q)$ ,  $T_i$  needs to read value of Q that was already overwritten

- ▶ read operation rejected and  $T_i$  rolled back

If  $TS(T_i) \geq W\text{-timestamp}(Q)$

- ▶ read executed, R-timestamp(Q) set to  $\max(R\text{-timestamp}(Q), TS(T_i))$





# Timestamp-ordering Protocol

Suppose  $T_i$  executes  $\text{write}(Q)$

If  $\text{TS}(T_i) < \text{R-timestamp}(Q)$ , value  $Q$  produced by  $T_i$  was needed previously and  $T_i$  assumed it would never be produced

- ▶ **Write** operation rejected,  $T_i$  rolled back

If  $\text{TS}(T_i) < \text{W-timestamp}(Q)$ ,  $T_i$  attempting to write obsolete value of  $Q$

- ▶ **Write** operation rejected and  $T_i$  rolled back

Otherwise, **write** executed

Any rolled back transaction  $T_i$  is assigned new timestamp and restarted

Algorithm ensures conflict serializability and freedom from deadlock







# Schedule Possible Under Timestamp Protocol

$T_2$	$T_3$
read( $B$ )	read( $B$ ) write( $B$ )
read( $A$ )	read( $A$ ) write( $A$ )

