

# Chapter 13: I/O Systems

---





# Chapter 13: I/O Systems

---

Overview

I/O Hardware

Application I/O Interface

Kernel I/O Subsystem

Transforming I/O Requests to Hardware Operations

STREAMS

Performance





# Objectives

---

Explore the structure of an operating system's I/O subsystem

Discuss the principles of I/O hardware and its complexity

Provide details of the performance aspects of I/O hardware and software





# Overview

---

I/O management is a major component of operating system design and operation

- Important aspect of computer operation

- I/O devices vary greatly

- Various methods to control them

- Performance management

- New types of devices frequent

Ports, busses, device controllers connect to various devices

**Device drivers** encapsulate device details

- Present uniform device-access interface to I/O subsystem





# I/O Hardware

Incredible variety of I/O devices

Storage

Transmission

Human-interface

Common concepts – signals from I/O devices interface with computer

**Port** – connection point for device

**Bus** - **daisy chain** or shared direct access

- ▶ **PCI** bus common in PCs and servers, PCI Express (**PCIe**)
- ▶ **expansion bus** connects relatively slow devices

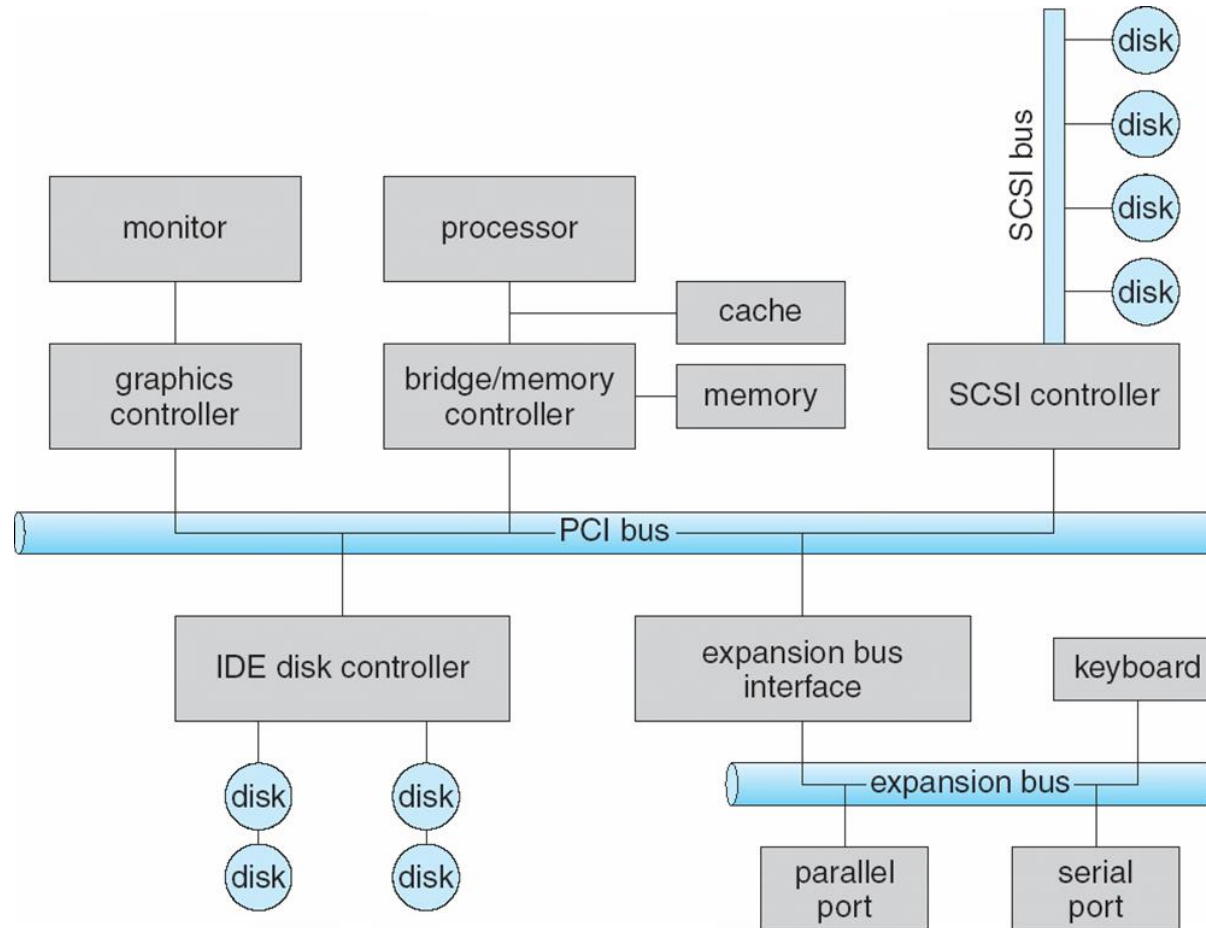
**Controller (host adapter)** – electronics that operate port, bus, device

- ▶ Sometimes integrated
- ▶ Sometimes separate circuit board (host adapter)
- ▶ Contains processor, microcode, private memory, bus controller, etc
  - Some talk to per-device controller with bus controller, microcode, memory, etc





# A Typical PC Bus Structure





# I/O Hardware (Cont.)

---

I/O instructions control devices

Devices usually have registers where device driver places commands, addresses, and data to write, or read data from registers after command execution

Data-in register, data-out register, status register, control register

Typically 1-4 bytes, or FIFO buffer

Devices have addresses, used by

Direct I/O instructions

## Memory-mapped I/O

- ▶ Device data and command registers mapped to processor address space
- ▶ Especially for large address spaces (graphics)





# Device I/O Port Locations on PCs (partial)

I/O address range (hexadecimal)	device
000–00F	DMA controller
020–021	interrupt controller
040–043	timer
200–20F	game controller
2F8–2FF	serial port (secondary)
320–32F	hard-disk controller
378–37F	parallel port
3D0–3DF	graphics controller
3F0–3F7	diskette-drive controller
3F8–3FF	serial port (primary)







# Polling

For each byte of I/O

1. Read busy bit from status register until 0
2. Host sets read or write bit and if write copies data into data-out register
3. Host sets command-ready bit
4. Controller sets busy bit, executes transfer
5. Controller clears busy bit, error bit, command-ready bit when transfer done

Step 1 is **busy-wait** cycle to wait for I/O from device

Reasonable if device is fast

But inefficient if device slow

CPU switches to other tasks?

- ▶ But if miss a cycle data overwritten / lost





# Interrupts

---

Polling can happen in 3 instruction cycles

- Read status, logical-and to extract status bit, branch if not zero

- How to be more efficient if non-zero infrequently?

CPU **Interrupt-request line** triggered by I/O device

- Checked by processor after each instruction

**Interrupt handler** receives interrupts

- Maskable** to ignore or delay some interrupts

**Interrupt vector** to dispatch interrupt to correct handler

- Context switch at start and end

- Based on priority


- Some **nonmaskable**

- Interrupt chaining if more than one device at same interrupt number





# Interrupt-Driven I/O Cycle





# Intel Pentium Processor Event-Vector Table

vector number	description
0	divide error
1	debug exception
2	null interrupt
3	breakpoint
4	INTO-detected overflow
5	bound range exception
6	invalid opcode
7	device not available
8	double fault
9	coprocessor segment overrun (reserved)
10	invalid task state segment
11	segment not present
12	stack fault
13	general protection
14	page fault
15	(Intel reserved, do not use)
16	floating-point error
17	alignment check
18	machine check
19–31	(Intel reserved, do not use)
32–255	maskable interrupts





# Interrupts (Cont.)

---

Interrupt mechanism also used for **exceptions**

- Terminate process, crash system due to hardware error

- Page fault executes when memory access error

- System call executes via **trap** to trigger kernel to execute request

- Multi-CPU systems can process interrupts concurrently

- If operating system designed to handle it

- Used for time-sensitive processing, frequent, must be fast





# Direct Memory Access

Used to avoid **programmed I/O** (one byte at a time) for large data movement

Requires **DMA** controller

Bypasses CPU to transfer data directly between I/O device and memory

OS writes DMA command block into memory

- Source and destination addresses

- Read or write mode

- Count of bytes

- Writes location of command block to DMA controller

- Bus mastering of DMA controller – grabs bus from CPU

  - ▶ **Cycle stealing** from CPU but still much more efficient

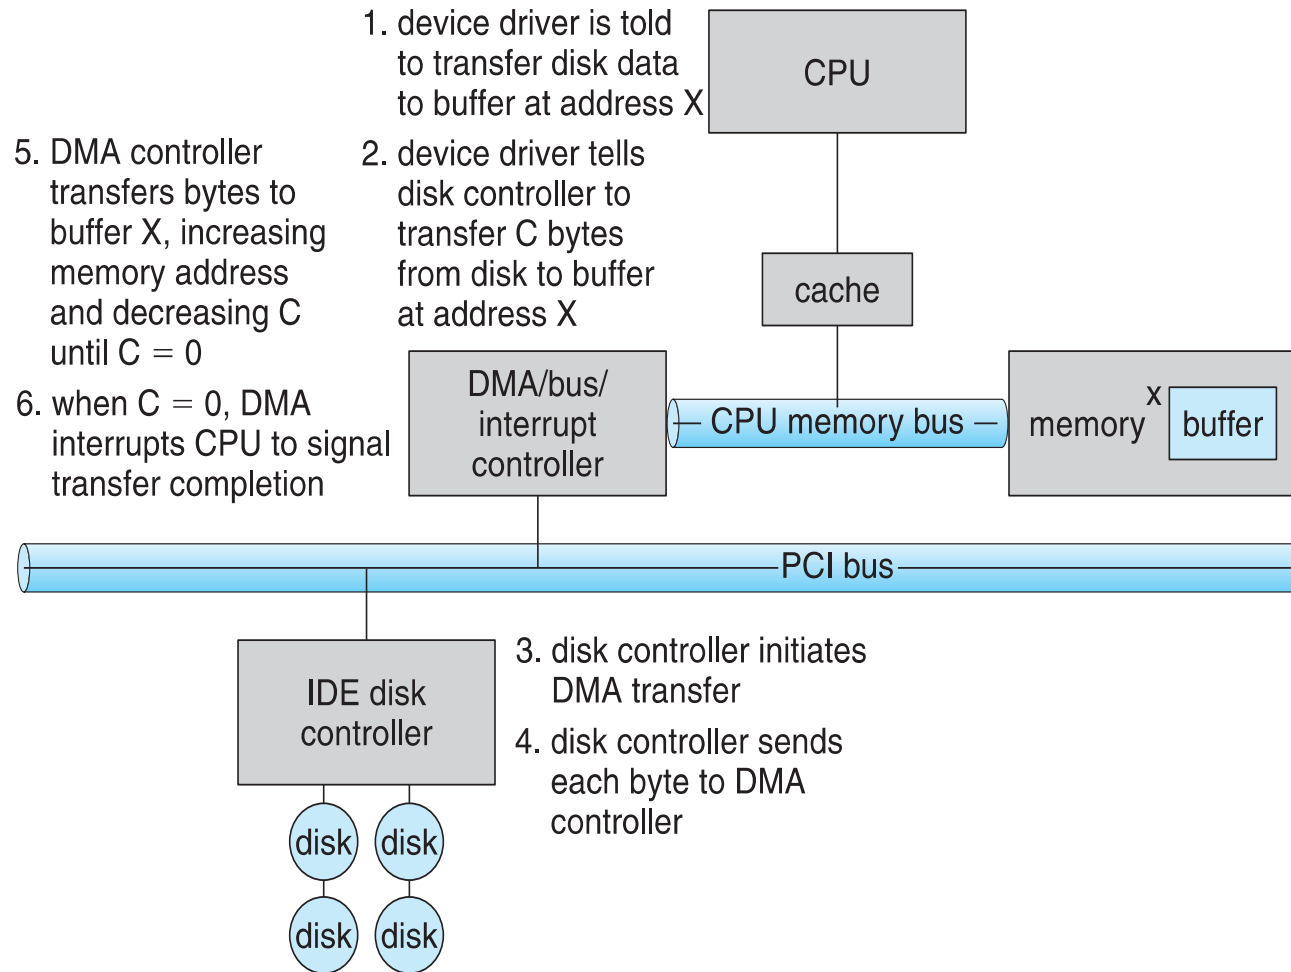
- When done, interrupts to signal completion

Version that is aware of virtual addresses can be even more efficient - **DVMA**





# Six Step Process to Perform DMA Transfer





# Application I/O Interface

---

I/O system calls encapsulate device behaviors in generic classes

Device-driver layer hides differences among I/O controllers from kernel

New devices talking already-implemented protocols need no extra work

Each OS has its own I/O subsystem structures and device driver frameworks

Devices vary in many dimensions

**Character-stream** or **block**

**Sequential** or **random-access**

**Synchronous** or **asynchronous** (or both)

**Sharable** or **dedicated**

**Speed of operation**

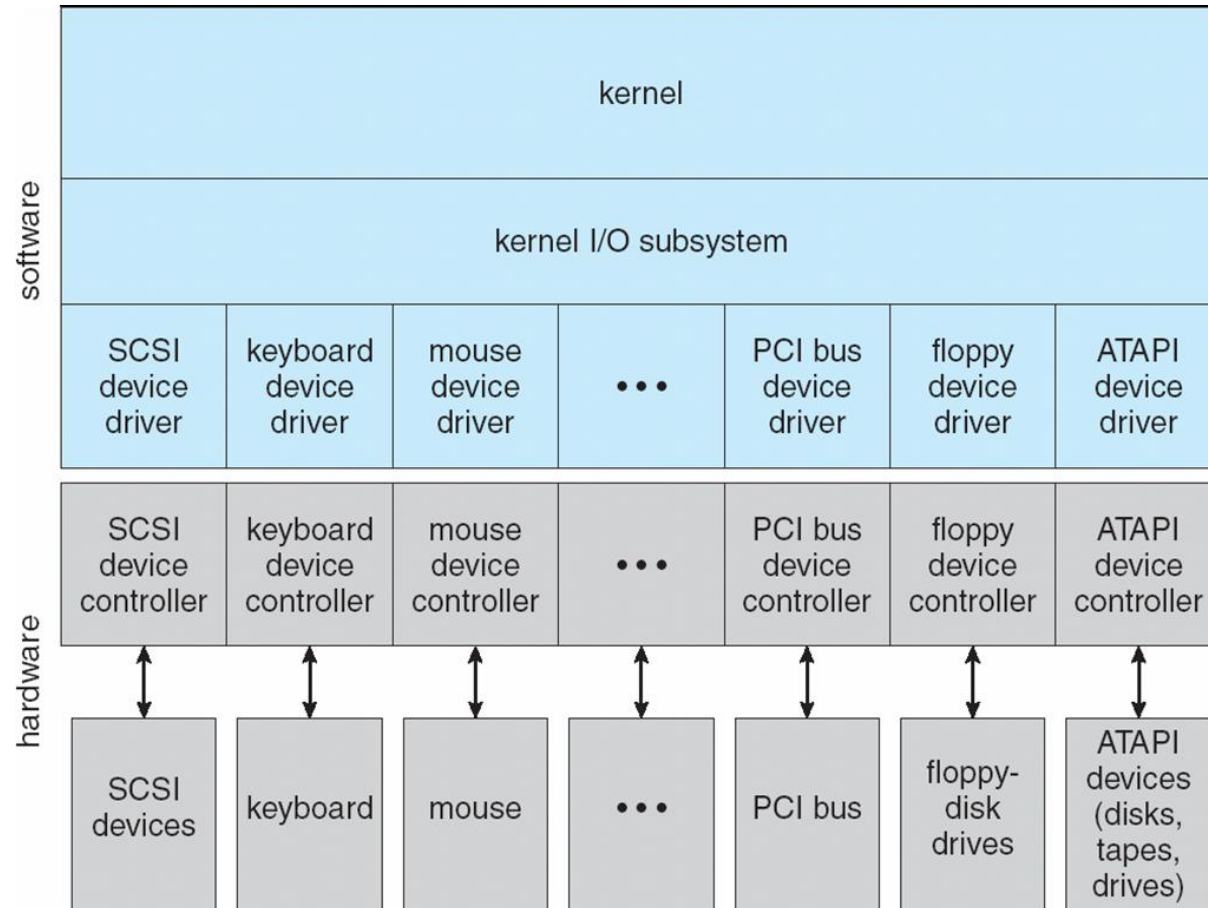
**read-write, read only, or write only**







# A Kernel I/O Structure





# Characteristics of I/O Devices

aspect	variation	example
data-transfer mode	character block	terminal disk
access method	sequential random	modem CD-ROM
transfer schedule	synchronous asynchronous	tape keyboard
sharing	dedicated sharable	tape keyboard
device speed	latency seek time transfer rate delay between operations	
I/O direction	read only write only read–write	CD-ROM graphics controller disk





# Characteristics of I/O Devices (Cont.)

---

Subtleties of devices handled by device drivers

Broadly I/O devices can be grouped by the OS into

- Block I/O

- Character I/O (Stream)

- Memory-mapped file access

- Network sockets

For direct manipulation of I/O device specific characteristics, usually an escape / back door

- Unix `ioctl()` call to send arbitrary bits to a device control register and data to device data register





# Block and Character Devices

---

Block devices include disk drives

Commands include read, write, seek

**Raw I/O**, **direct I/O**, or file-system access

Memory-mapped file access possible

- ▶ File mapped to virtual memory and clusters brought via demand paging

DMA

Character devices include keyboards, mice, serial ports

Commands include `get()`, `put()`

Libraries layered on top allow line editing





# Network Devices

---

Varying enough from block and character to have own interface

Linux, Unix, Windows and many others include **socket** interface

- Separates network protocol from network operation

- Includes **select()** functionality

Approaches vary widely (pipes, FIFOs, streams, queues, mailboxes)





# Clocks and Timers

---

Provide current time, elapsed time, timer

Normal resolution about 1/60 second

Some systems provide higher-resolution timers

**Programmable interval timer** used for timings, periodic interrupts

`ioctl()` (on UNIX) covers odd aspects of I/O such as clocks and timers





# Nonblocking and Asynchronous I/O

---

**Blocking** - process suspended until I/O completed

Easy to use and understand

Insufficient for some needs

**Nonblocking** - I/O call returns as much as available

User interface, data copy (buffered I/O)

Implemented via multi-threading

Returns quickly with count of bytes read or written

`select()` to find if data ready then `read()` or `write()` to transfer

**Asynchronous** - process runs while I/O executes

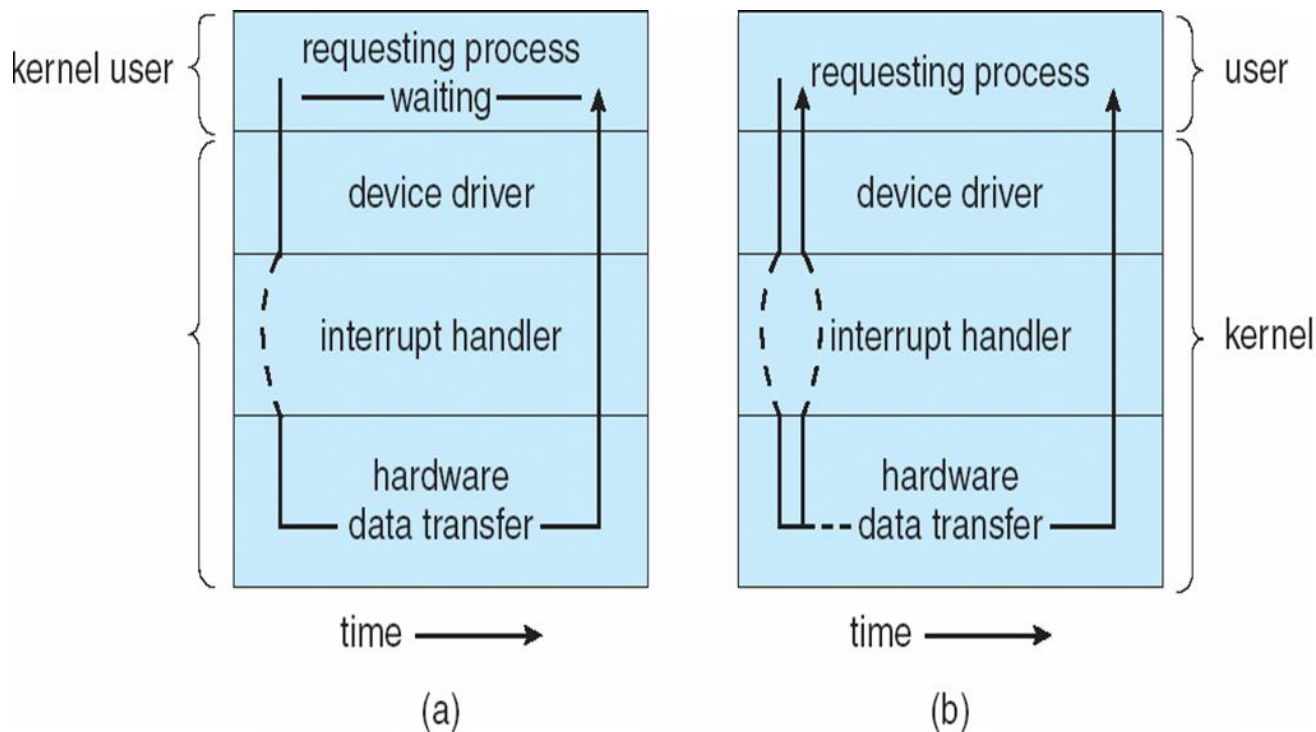
Difficult to use

I/O subsystem signals process when I/O completed





# Two I/O Methods



Synchronous

Asynchronous







# Vectored I/O

---

**Vectored I/O** allows one system call to perform multiple I/O operations

For example, Unix **readve()** accepts a vector of multiple buffers to read into or write from

This scatter-gather method better than multiple individual I/O calls

Decreases context switching and system call overhead

Some versions provide atomicity

- ▶ Avoid for example worry about multiple threads changing data as reads / writes occurring





# Kernel I/O Subsystem

---

## Scheduling

Some I/O request ordering via per-device queue

Some OSs try fairness

Some implement Quality Of Service (i.e. IPQOS)

**Buffering** - store data in memory while transferring between devices

To cope with device speed mismatch

To cope with device transfer size mismatch

To maintain “copy semantics”

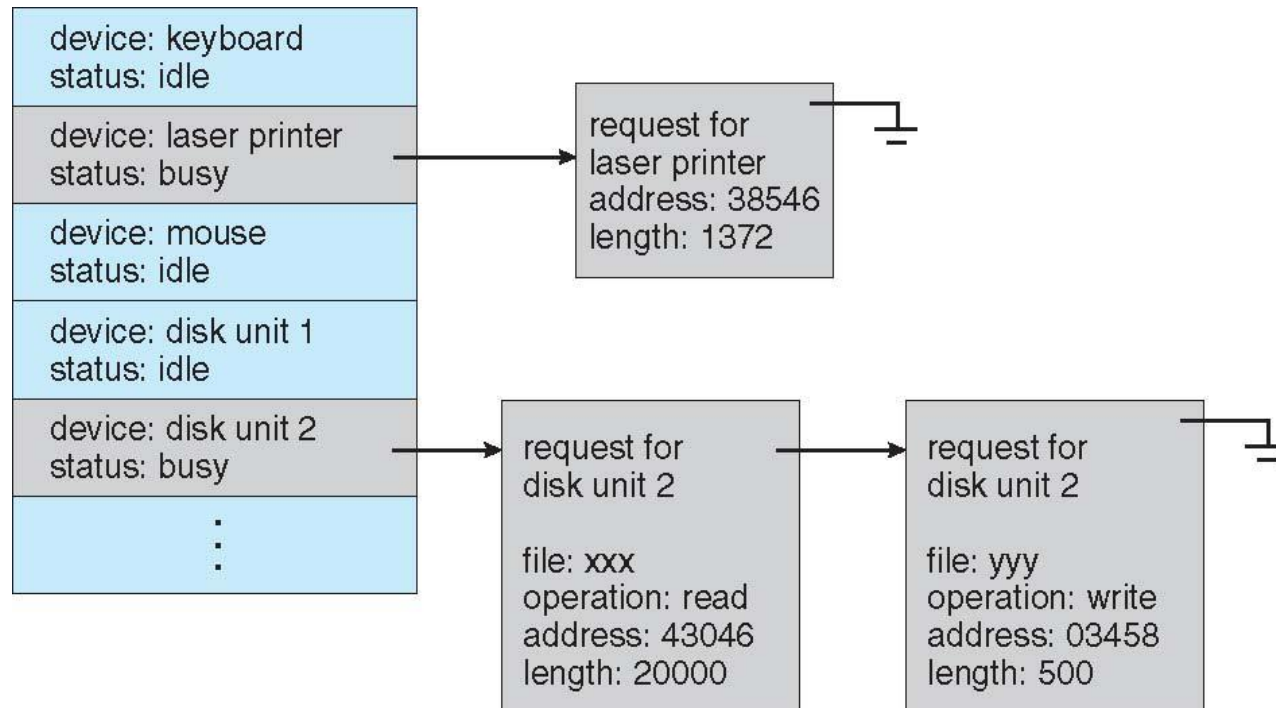
**Double buffering** – two copies of the data

- ▶ Kernel and user
- ▶ Varying sizes
- ▶ Full / being processed and not-full / being used
- ▶ Copy-on-write can be used for efficiency in some cases






# Device-status Table





# Sun Enterprise 6000 Device-Transfer Rates





# Kernel I/O Subsystem

---

**Caching** - faster device holding copy of data

- Always just a copy

- Key to performance

- Sometimes combined with buffering

**Spooling** - hold output for a device

- If device can serve only one request at a time

- i.e., Printing

**Device reservation** - provides exclusive access to a device

- System calls for allocation and de-allocation

- Watch out for deadlock





# Error Handling

---

OS can recover from disk read, device unavailable, transient write failures

Retry a read or write, for example

Some systems more advanced – Solaris FMA, AIX

- ▶ Track error frequencies, stop using device with increasing frequency of retry-able errors

Most return an error number or code when I/O request fails

System error logs hold problem reports





# I/O Protection

---

User process may accidentally or purposefully attempt to disrupt normal operation via illegal I/O instructions

All I/O instructions defined to be privileged

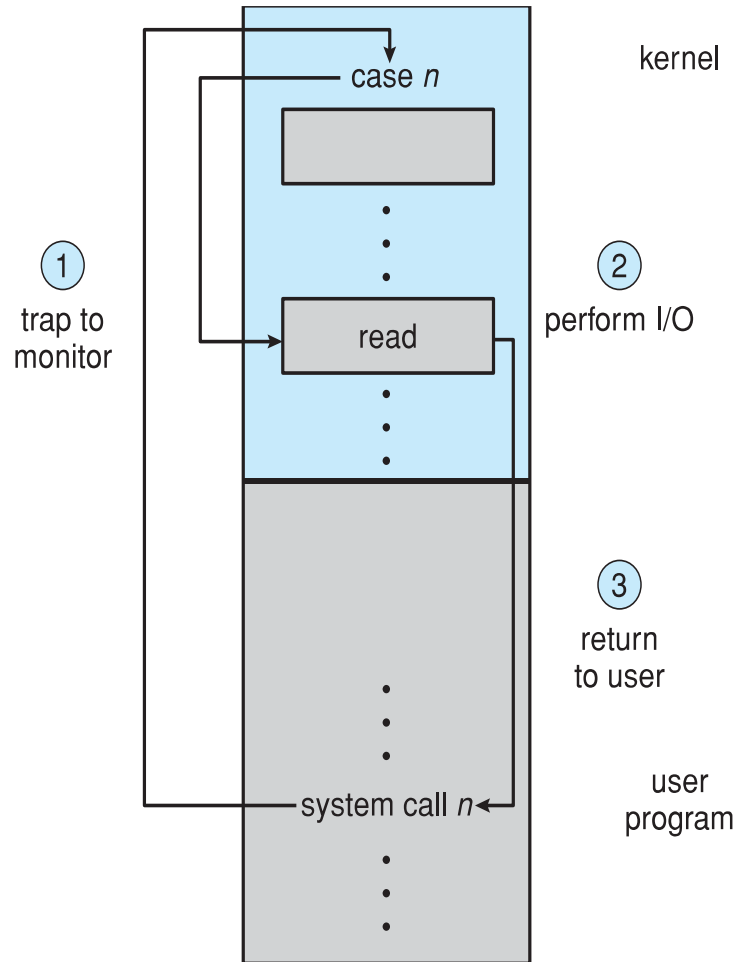
I/O must be performed via system calls

- ▶ Memory-mapped and I/O port memory locations must be protected too





# Use of a System Call to Perform I/O







# Kernel Data Structures

---

Kernel keeps state info for I/O components, including open file tables, network connections, character device state

Many, many complex data structures to track buffers, memory allocation, “dirty” blocks

Some use object-oriented methods and message passing to implement I/O

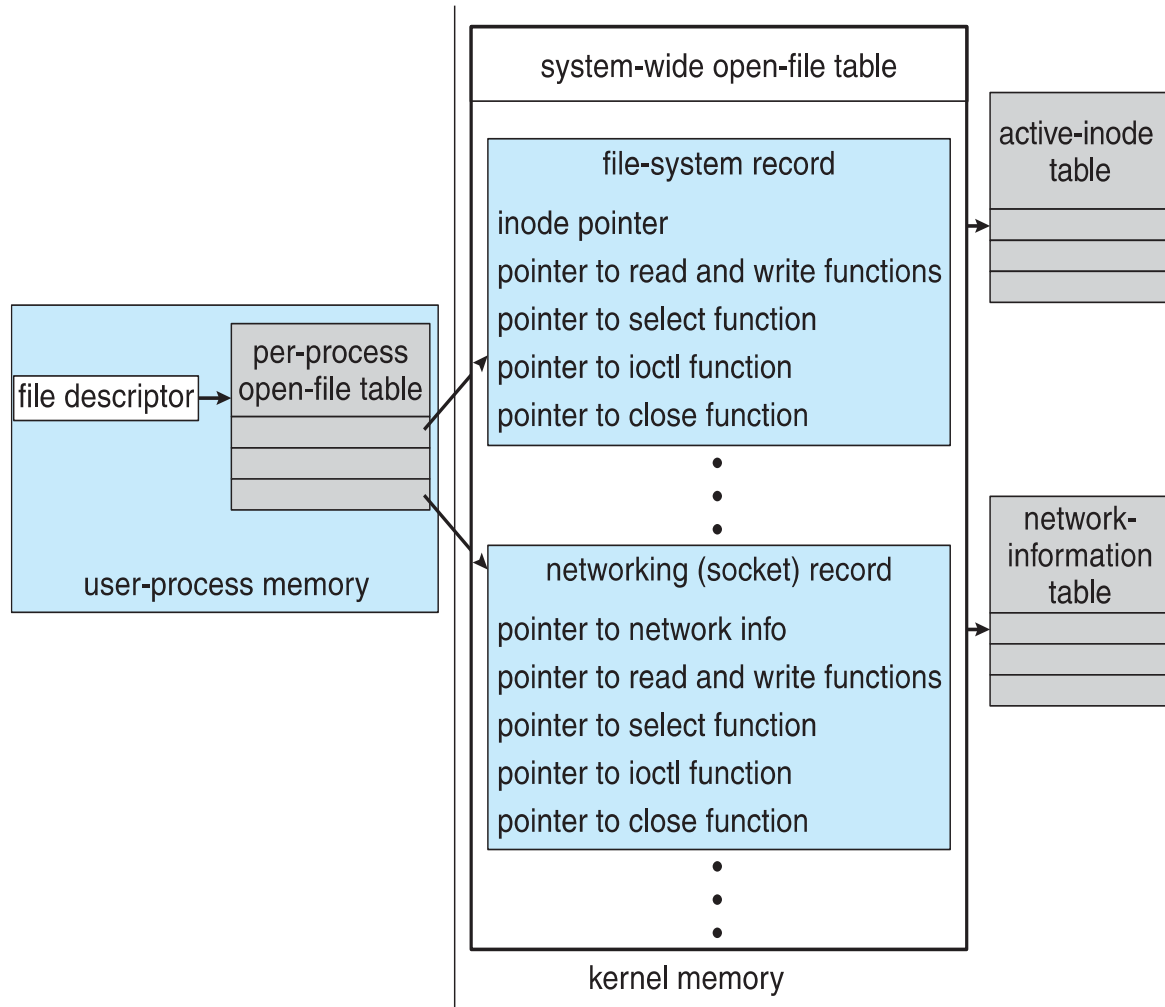
Windows uses message passing

- ▶ Message with I/O information passed from user mode into kernel
- ▶ Message modified as it flows through to device driver and back to process
- ▶ Pros / cons?





# UNIX I/O Kernel Structure





# Power Management

---

Not strictly domain of I/O, but much is I/O related

Computers and devices use electricity, generate heat, frequently require cooling

OSes can help manage and improve use

Cloud computing environments move virtual machines between servers

- ▶ Can end up evacuating whole systems and shutting them down

Mobile computing has power management as first class OS aspect





# Power Management (Cont.)

---

For example, Android implements

Component-level power management

- ▶ Understands relationship between components
- ▶ Build device tree representing physical device topology
- ▶ System bus -> I/O subsystem -> {flash, USB storage}
- ▶ Device driver tracks state of device, whether in use
- ▶ Unused component – turn it off
- ▶ All devices in tree branch unused – turn off branch

Wake locks – like other locks but prevent sleep of device when lock is held

Power collapse – put a device into very deep sleep

- ▶ Marginal power use
- ▶ Only awake enough to respond to external stimuli (button press, incoming call)





# I/O Requests to Hardware Operations

---

Consider reading a file from disk for a process:

- Determine device holding file

- Translate name to device representation

- Physically read data from disk into buffer

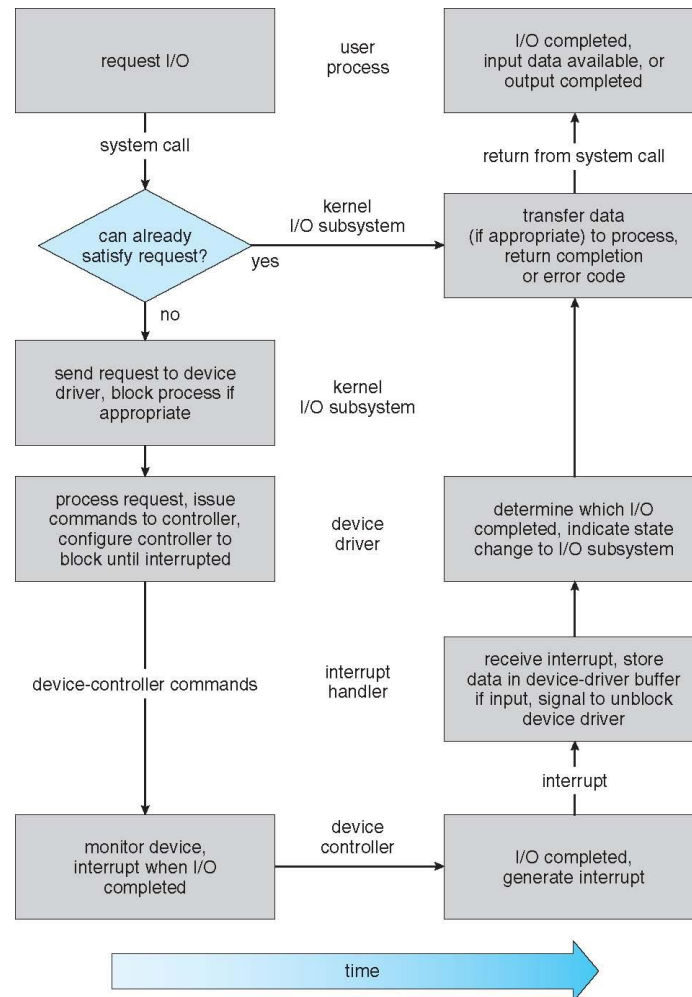
- Make data available to requesting process

- Return control to process





# Life Cycle of An I/O Request





# STREAMS

---

**STREAM** – a full-duplex communication channel between a user-level process and a device in Unix System V and beyond

A STREAM consists of:

- STREAM head interfaces with the user process

- driver end interfaces with the device

- zero or more STREAM modules between them

Each module contains a **read queue** and a **write queue**

Message passing is used to communicate between queues

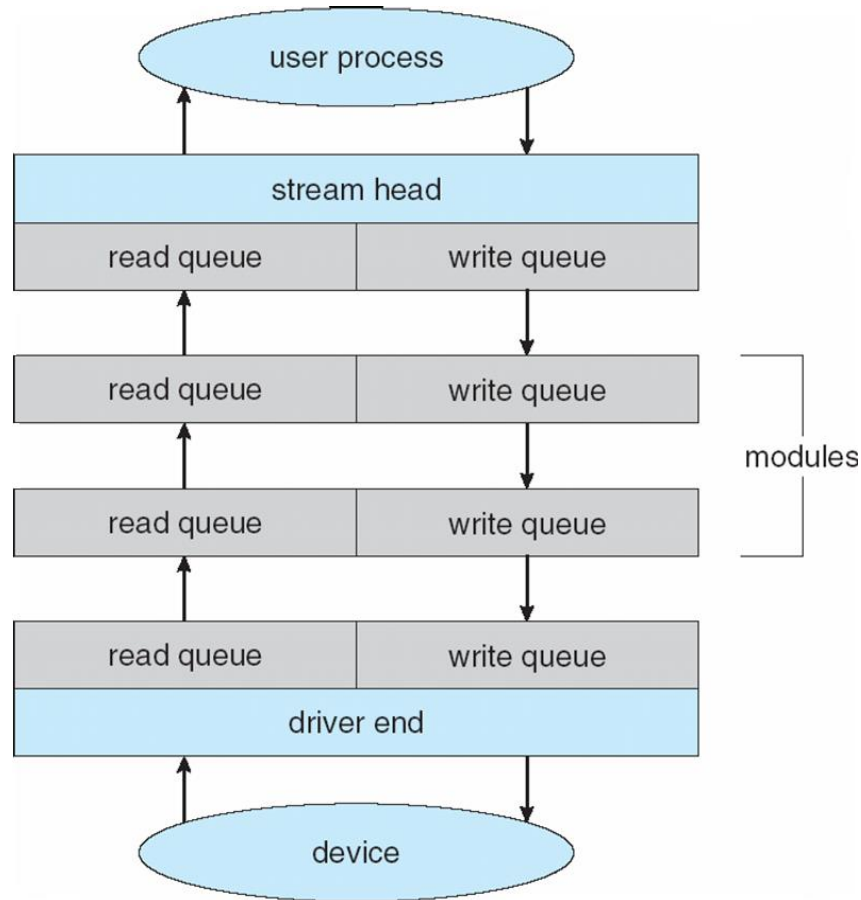
- Flow control** option to indicate available or busy

Asynchronous internally, synchronous where user process communicates with stream head





# The STREAMS Structure







# Performance

---

I/O a major factor in system performance:

- Demands CPU to execute device driver, kernel I/O code

- Context switches due to interrupts

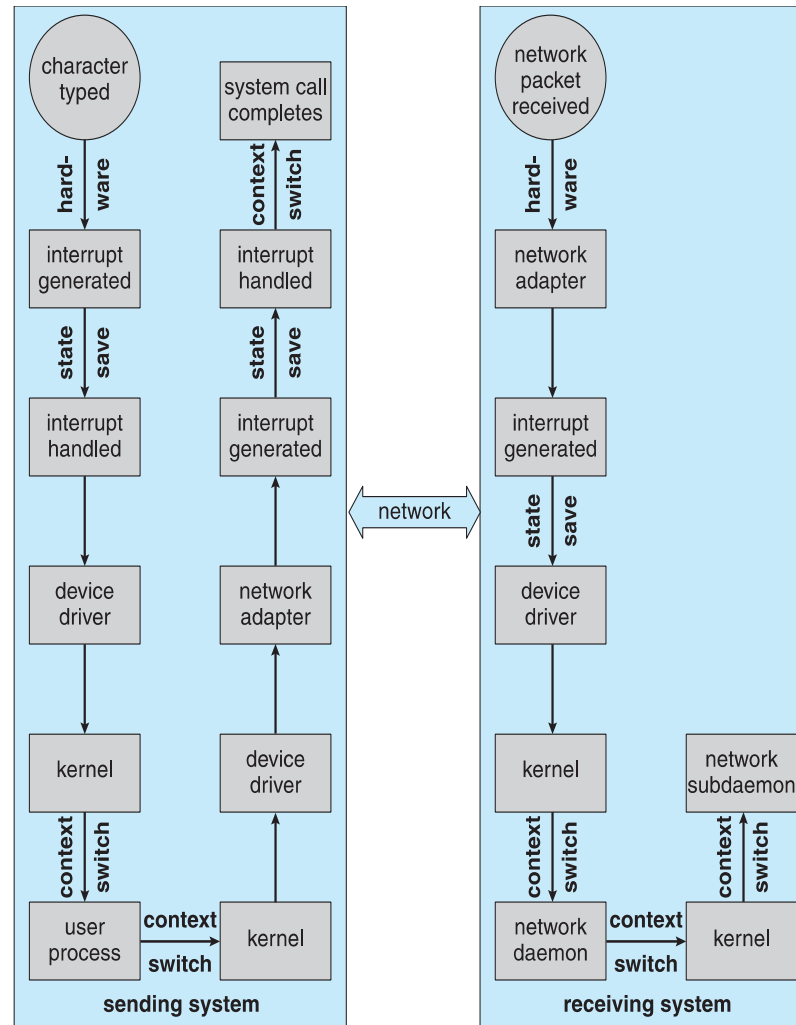
- Data copying

- Network traffic especially stressful





# Intercomputer Communications





# Improving Performance

---

Reduce number of context switches

Reduce data copying

Reduce interrupts by using large transfers, smart controllers, polling

Use DMA

Use smarter hardware devices

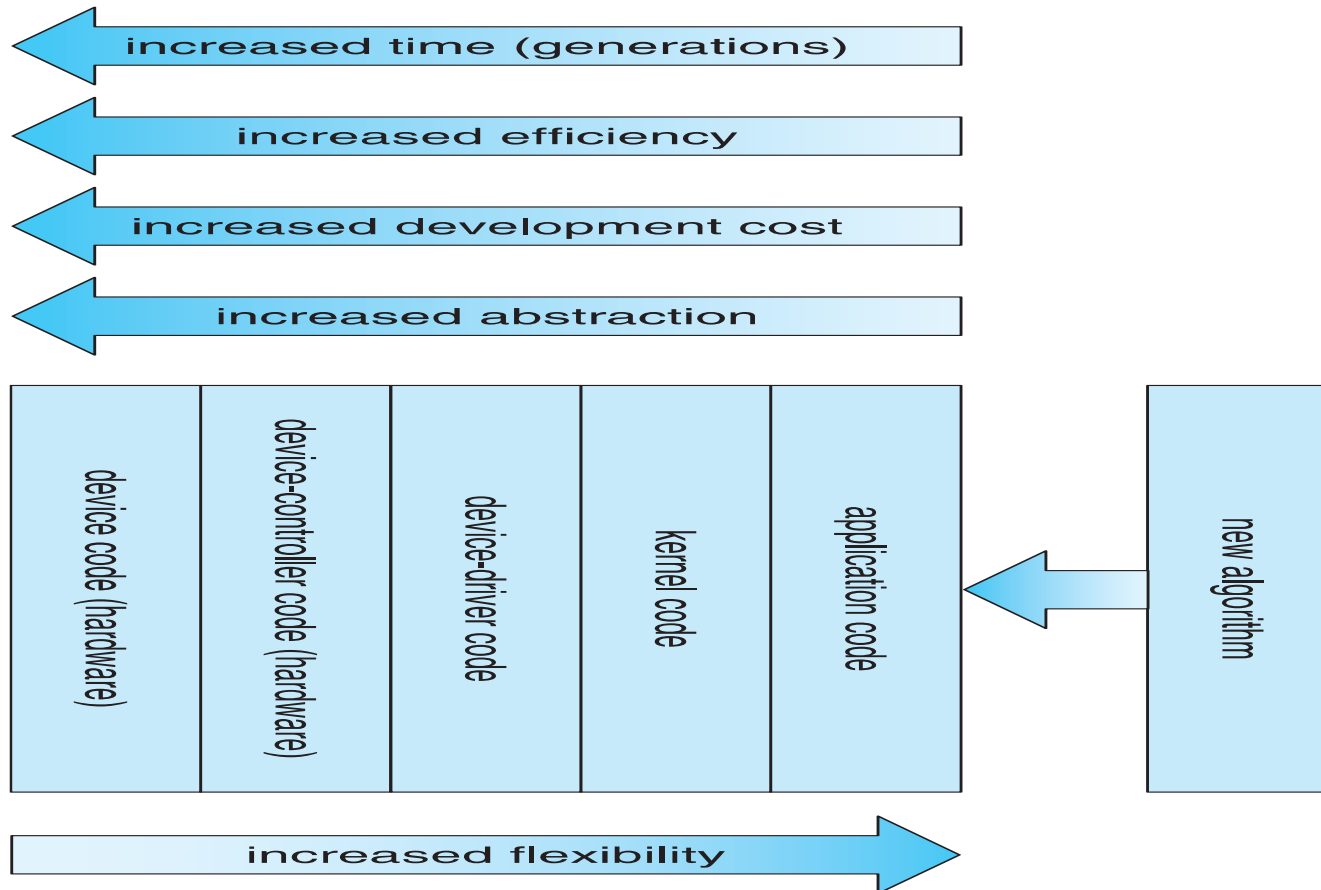
Balance CPU, memory, bus, and I/O performance for highest throughput

Move user-mode processes / daemons to kernel threads





# Device-Functionality Progression



# End of Chapter 13

---

