

Chapter 17:

Distributed Systems





Chapter 17: Distributed Systems

Advantages of Distributed Systems

Types of Network-Based Operating Systems

Network Structure

Communication Structure

Communication Protocols

An Example: TCP/IP

Robustness

Design Issues

Distributed File System





Chapter Objectives

To provide a high-level overview of distributed systems and the networks that interconnect them

To discuss the general structure of distributed operating systems

To explain general communication structure and communication protocols


To describe issues concerning the design of distributed systems





Overview

- n **Distributed system** is collection of loosely coupled processors interconnected by a communications network
- n Processors variously called **nodes**, **computers**, **machines**, **hosts**
 - | **Site** is location of the processor
 - | Generally a **server** has a resource a **client** node at a different site wants to use





Reasons for Distributed Systems

Reasons for distributed systems

Resource sharing

- ▶ Sharing and printing files at remote sites
- ▶ Processing information in a distributed database
- ▶ Using remote specialized hardware devices

Computation speedup – load sharing or job migration

Reliability – detect and recover from site failure, function transfer, reintegrate failed site

Communication – message passing

- ▶ All higher-level functions of a standalone system can be expanded to encompass a distributed system

Computers can be downsized, more flexibility, better user interfaces and easier maintenance by moving from large system to multiple smaller systems performing distributed computing





Types of Distributed Operating Systems

Network Operating Systems

Distributed Operating Systems





Network-Operating Systems

Users are aware of multiplicity of machines

Access to resources of various machines is done explicitly by:

- Remote logging into the appropriate remote machine (telnet, ssh)

- Remote Desktop (Microsoft Windows)

- Transferring data from remote machines to local machines, via the File Transfer Protocol (FTP) mechanism

Users must change paradigms – establish a **session**, give network-based commands

- More difficult for users





Distributed-Operating Systems

Users not aware of multiplicity of machines

Access to remote resources similar to access to local resources

Data Migration – transfer data by transferring entire file, or transferring only those portions of the file necessary for the immediate task

Computation Migration – transfer the computation, rather than the data, across the system

Via remote procedure calls (RPCs)
or via messaging system





Distributed-Operating Systems (Cont.)

Process Migration – execute an entire process, or parts of it, at different sites

Load balancing – distribute processes across network to even the workload

Computation speedup – subprocesses can run concurrently on different sites

Hardware preference – process execution may require specialized processor

Software preference – required software may be available at only a particular site

Data access – run process remotely, rather than transfer all data locally

Consider the World Wide Web





Network Structure

Local-Area Network (LAN) – designed to cover small geographical area

Multiple topologies like star or ring

Speeds from 1Mb per second (Appletalk, bluetooth) to 40 Gbps for fastest Ethernet over twisted pair copper or optical fibre

Consists of multiple computers (mainframes through mobile devices), peripherals (printers, storage arrays), routers (specialized network communication processors) providing access to other networks


Ethernet most common way to construct LANs

- ▶ Multiaccess bus-based
- ▶ Defined by standard IEEE 802.3


Wireless spectrum (**WiFi**) increasingly used for networking

- ▶ I.e. IEEE 802.11g standard implemented at 54 Mbps





Local-area Network





Network Types (Cont.)

Wide-Area Network (WAN) – links geographically separated sites

Point-to-point connections over long-haul lines (often leased from a phone company)

- ▶ Implemented via **connection processors** known as **routers**

Internet WAN enables hosts world wide to communicate

- ▶ Hosts differ in all dimensions but WAN allows communications

Speeds

- ▶ T1 link is 1.544 Megabits per second
- ▶ T3 is $28 \times T1s = 45 \text{ Mbps}$
- ▶ OC-12 is 622 Mbps


WANs and LANs interconnect, similar to cell phone network:

- ▶ Cell phones use radio waves to cell towers
- ▶ Towers connect to other towers and hubs





Communication Processors in a Wide-Area Network





Communication Structure

The design of a communication network must address four basic issues:

Naming and name resolution - How do two processes locate each other to communicate?

Routing strategies - How are messages sent through the network?

Connection strategies - How do two processes send a sequence of messages?

Contention - The network is a shared resource, so how do we resolve conflicting demands for its use?





Naming and Name Resolution

Name systems in the network

Address messages with the process-id

Identify processes on remote systems by

<**host-name**, **identifier**> pair

Domain name system (DNS) – specifies the naming structure of the hosts, as well as name to address **resolution** (Internet)





Routing Strategies

Fixed routing - A path from A to B is specified in advance; path changes only if a hardware failure disables it

Since the shortest path is usually chosen, communication costs are minimized

Fixed routing cannot adapt to load changes

Ensures that messages will be delivered in the order in which they were sent

Virtual routing - A path from A to B is fixed for the duration of one session. Different sessions involving messages from A to B may have different paths

Partial remedy to adapting to load changes

Ensures that messages will be delivered in the order in which they were sent





Routing Strategies (Cont.)

Dynamic routing - The path used to send a message from site *A* to site *B* is chosen only when a message is sent

Usually a site sends a message to another site on the link least used at that particular time

Adapts to load changes by avoiding routing messages on heavily used path

Messages may arrive out of order

- ▶ This problem can be remedied by appending a sequence number to each message

Most complex to set up

Tradeoffs mean all methods are used

UNIX provides ability to mix fixed and dynamic

Hosts may have fixed routes and **gateways** connecting networks together may have dynamic routes





Routing Strategies (Cont.)

Router is communications processor responsible for routing messages

Must have at least 2 network connections

Maybe special purpose or just function running on host

Checks its tables to determine where destination host is, where to send messages

Static routing – table only changed manually

Dynamic routing – table changed via **routing protocol**





Routing Strategies (Cont.)

More recently, routing managed by intelligent software more intelligently than routing protocols

OpenFlow is device-independent, allowing developers to introduce network efficiencies by decoupling data-routing decisions from underlying network devices

Messages vary in length – simplified design breaks them into **packets** (or **frames**, or **datagrams**)

Connectionless message is just one packet

Otherwise need a connection to get a multi-packet message from source to destination





Connection Strategies

Circuit switching - A permanent physical link is established for the duration of the communication (i.e., telephone system)

Message switching - A temporary link is established for the duration of one message transfer (i.e., post-office mailing system)

Packet switching - Messages of variable length are divided into fixed-length packets which are sent to the destination

Each packet may take a different path through the network

The packets must be reassembled into messages as they arrive

Circuit switching requires setup time, but incurs less overhead for shipping each message, and may waste network bandwidth

Message and packet switching require less setup time, but incur more overhead per message





Communication Protocol

The communication network is partitioned into the following multiple layers:

Layer 1: Physical layer – handles the mechanical and electrical details of the physical transmission of a bit stream

Layer 2: Data-link layer – handles the *frames*, or fixed-length parts of packets, including any error detection and recovery that occurred in the physical layer

Layer 3: Network layer – provides connections and routes packets in the communication network, including handling the address of outgoing packets, decoding the address of incoming packets, and maintaining routing information for proper response to changing load levels





Communication Protocol (Cont.)

Layer 4: Transport layer – responsible for low-level network access and for message transfer between clients, including partitioning messages into packets, maintaining packet order, controlling flow, and generating physical addresses

Layer 5: Session layer – implements sessions, or process-to-process communications protocols

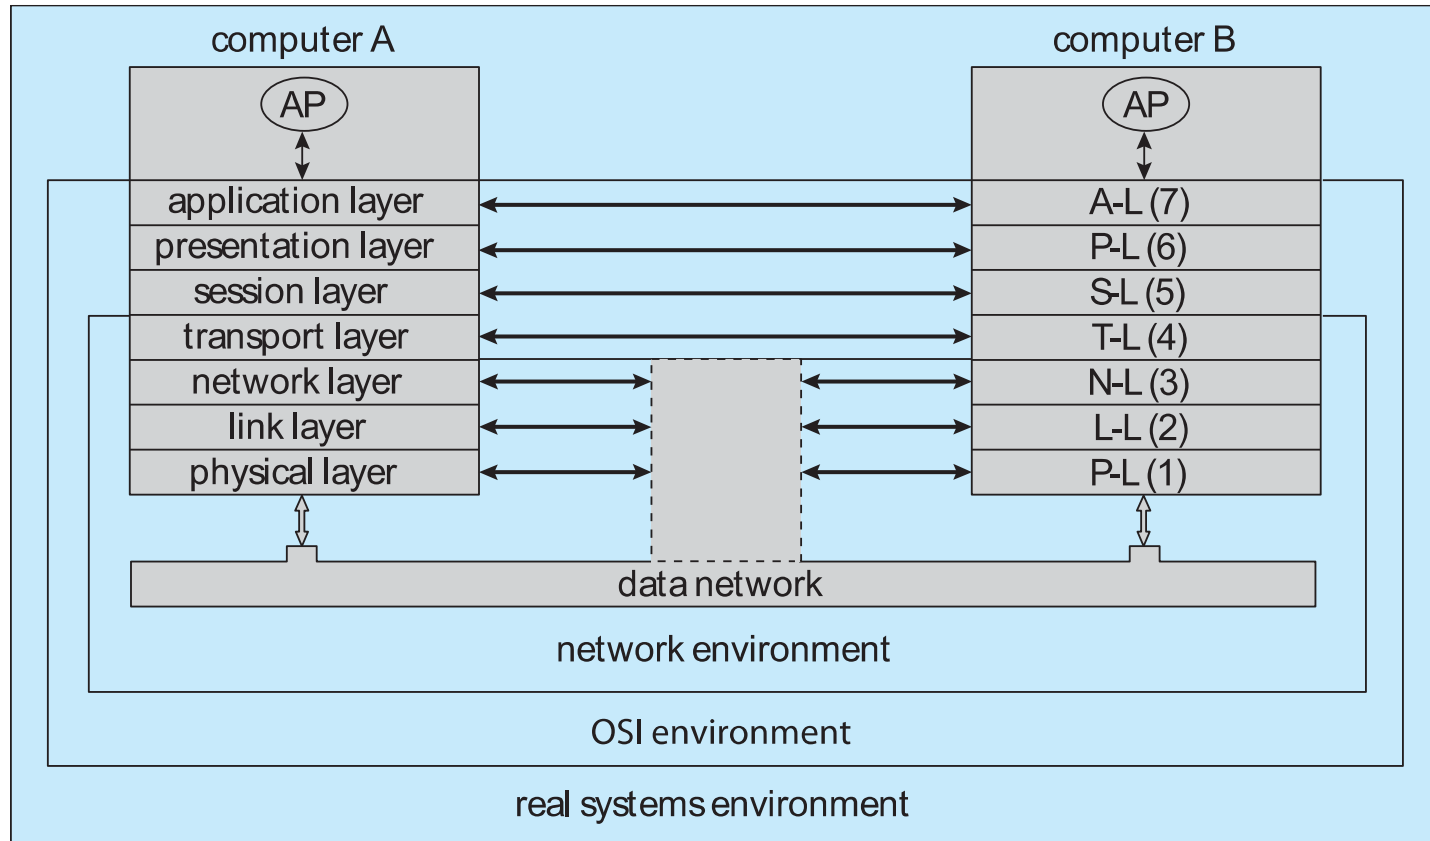
Layer 6: Presentation layer – resolves the differences in formats among the various sites in the network, including character conversions, and half duplex/full duplex (echoing)

Layer 7: Application layer – interacts directly with the users, deals with file transfer, remote-login protocols and electronic mail, as well as schemas for distributed databases



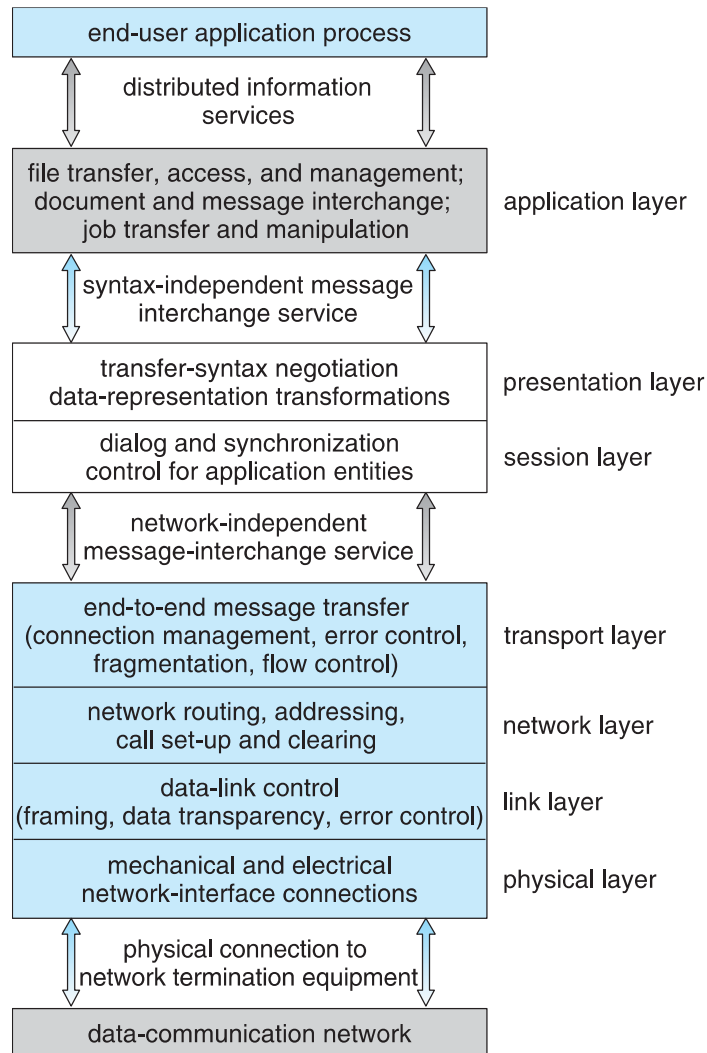


Communication Via ISO Network Model



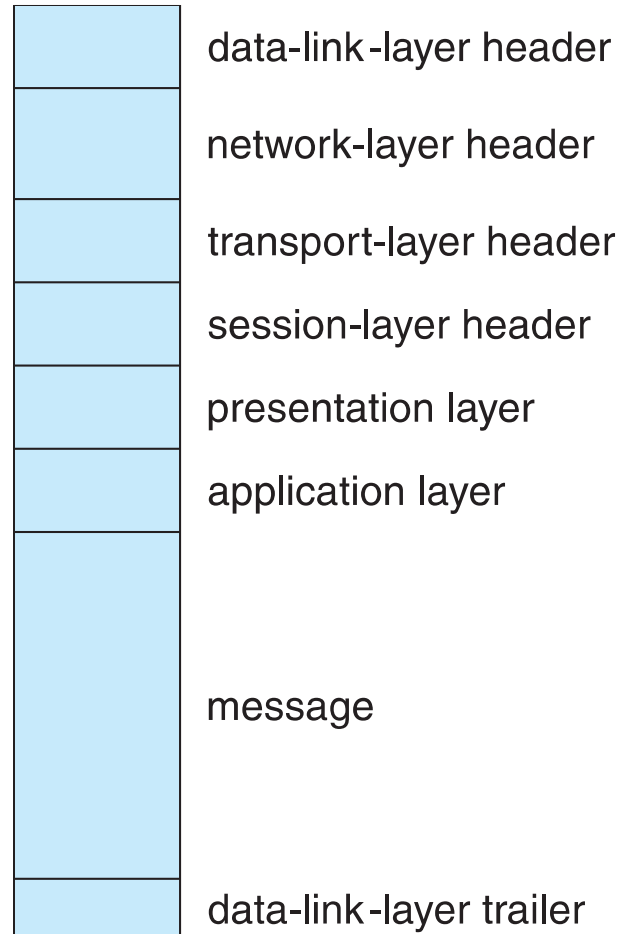


The ISO Protocol Layer



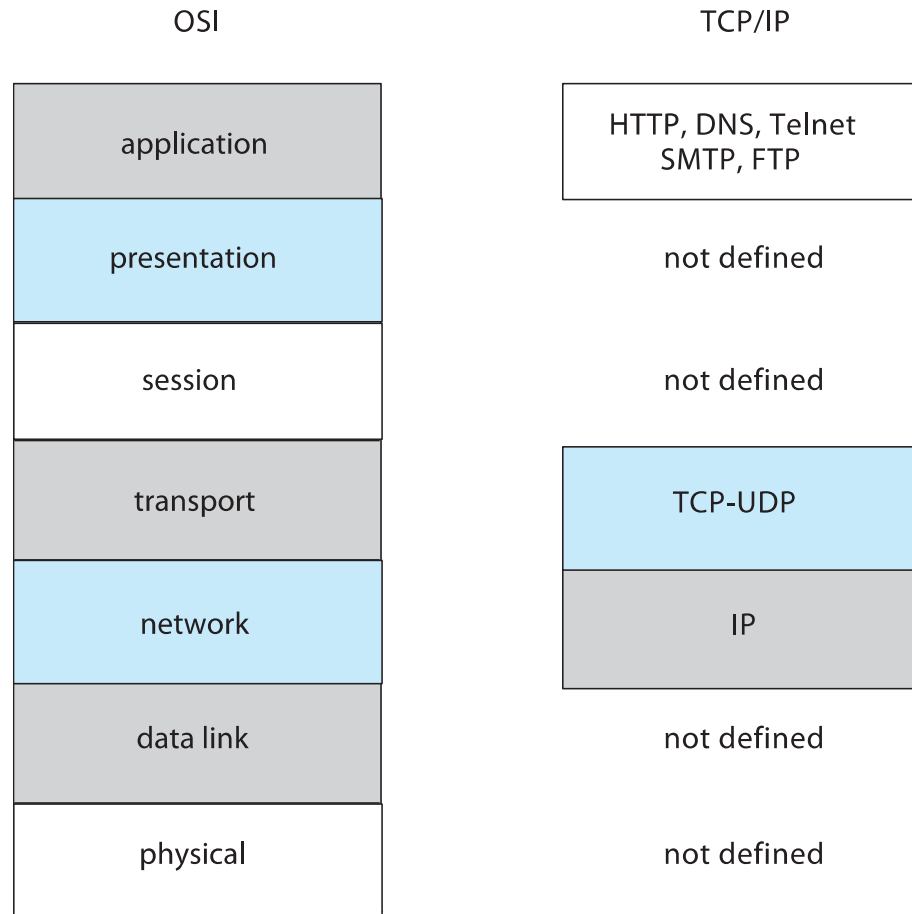


The ISO Network Message





The TCP/IP Protocol Layers





Example: TCP/IP

The transmission of a network packet between hosts on an Ethernet network

Every host has a unique IP address and a corresponding Ethernet **Media Access Control (MAC)** address

Communication requires both addresses

Domain Name Service (DNS) can be used to acquire IP addresses

Address Resolution Protocol (ARP) is used to map MAC addresses to IP addresses

Broadcast to all other systems on the Ethernet network

If the hosts are on the same network, ARP can be used

If the hosts are on different networks, the sending host will send the packet to a router which routes the packet to the destination network





An Ethernet Packet

bytes		
7	preamble—start of packet	each byte pattern 10101010
1	start of frame delimiter	pattern 10101011
2 or 6	destination address	Ethernet address or broadcast
2 or 6	source address	Ethernet address
2	length of data section	length in bytes
0–1500	data	message data
0–46	pad (optional)	message must be > 63 bytes long
4	frame checksum	for error detection





Robustness

Failure detection

Reconfiguration





Failure Detection

Detecting hardware failure is difficult

To detect a link failure, a **heartbeat** protocol can be used

Assume Site A and Site B have established a link

At fixed intervals, each site will exchange an *I-am-up* message indicating that they are up and running

If Site A does not receive a message within the fixed interval, it assumes either (a) the other site is not up or (b) the message was lost

Site A can now send an *Are-you-up?* message to Site B

If Site A does not receive a reply, it can repeat the message or try an alternate route to Site B





Failure Detection (Cont.)

If Site A does not ultimately receive a reply from Site B, it concludes some type of failure has occurred

Types of failures:

- Site B is down
- The direct link between A and B is down
- The alternate link from A to B is down
- The message has been lost

However, Site A cannot determine exactly **why** the failure has occurred





Reconfiguration

When Site A determines a failure has occurred, it must reconfigure the system:

1. If the link from A to B has failed, this must be broadcast to every site in the system
2. If a site has failed, every other site must also be notified indicating that the services offered by the failed site are no longer available

When the link or the site becomes available again, this information must again be broadcast to all other sites





Design Issues

Transparency – the distributed system should appear as a conventional, centralized system to the user

Fault tolerance – the distributed system should continue to function in the face of failure

Scalability – as demands increase, the system should easily accept the addition of new resources to accommodate the increased demand

Consider **Hadoop** open source programming framework for processing large datasets in distributed environments (based on Google search indexing)

Clusters – a collection of semi-autonomous machines that acts as a single system





Distributed File System

Distributed file system (DFS) – a distributed implementation of the classical time-sharing model of a file system, where multiple users share files and storage resources

A DFS manages set of dispersed storage devices

Overall storage space managed by a DFS is composed of different, remotely located, smaller storage spaces

There is usually a correspondence between constituent storage spaces and sets of files

Challenges include:

- Naming and Transparency

- Remote File Access





DFS Structure

Service – software entity running on one or more machines and providing a particular type of function to a priori unknown clients

Server – service software running on a single machine

Client – process that can invoke a service using a set of operations that forms its client interface

A client interface for a file service is formed by a set of primitive file operations (create, delete, read, write)

Client interface of a DFS should be transparent, i.e., not distinguish between local and remote files

Sometimes lower level **intermachine** interface need for cross-machine interaction





Naming and Transparency

Naming – mapping between logical and physical objects

Multilevel mapping – abstraction of a file that hides the details of how and where on the disk the file is actually stored

A **transparent** DFS hides the location where in the network the file is stored

For a file being **replicated** in several sites, the mapping returns a set of the locations of this file's replicas; both the existence of multiple copies and their location are hidden





Naming Structures

Location transparency – file name does not reveal the file's physical storage location

Location independence – file name does not need to be changed when the file's physical storage location changes





Naming Schemes — Three Main Approaches

Files named by combination of their host name and local name;
guarantees a unique system-wide name

Attach remote directories to local directories, giving the
appearance of a coherent directory tree; only previously mounted
remote directories can be accessed transparently

Total integration of the component file systems

- A single global name structure spans all the files in the system

- If a server is unavailable, some arbitrary set of directories on
different machines also becomes unavailable

In practice most DFSs use static, location-transparent mapping for
user-level names

- Some support file migration

- Hadoop supports file migration but without following POSIX
standards





Remote File Access

Remote-service mechanism is one transfer approach

Reduce network traffic by retaining recently accessed disk blocks in a cache, so that repeated accesses to the same information can be handled locally

If needed data not already cached, a copy of data is brought from the server to the user

Accesses are performed on the cached copy

Files identified with one master copy residing at the server machine, but copies of (parts of) the file are scattered in different caches

Cache-consistency problem – keeping the cached copies consistent with the master file

- ▶ Could be called **network virtual memory**





Cache Location – Disk vs. Main Memory

Advantages of disk caches

- More reliable

- Cached data kept on disk are still there during recovery and don't need to be fetched again

Advantages of main-memory caches:

- Permit workstations to be diskless

- Data can be accessed more quickly

- Performance speedup in bigger memories

- Server caches (used to speed up disk I/O) are in main memory regardless of where user caches are located; using main-memory caches on the user machine permits a single caching mechanism for servers and users





Cache Update Policy

Write-through – write data through to disk as soon as they are placed on any cache

Reliable, but poor performance

Delayed-write (write-back) – modifications written to the cache and then written through to the server later

Write accesses complete quickly; some data may be overwritten before they are written back, and so need never be written at all

Poor reliability; unwritten data will be lost whenever a user machine crashes

Variation – scan cache at regular intervals and flush blocks that have been modified since the last scan

Variation – **write-on-close**, writes data back to the server when the file is closed

- ▶ Best for files that are open for long periods and frequently modified





Consistency

Is locally cached copy of the data consistent with the master copy?

Client-initiated approach

- Client initiates a validity check

- Server checks whether the local data are consistent with the master copy

Server-initiated approach

- Server records, for each client, the (parts of) files it caches

- When server detects a potential inconsistency, it must react



End of Chapter 17

